

Why Learn About Op Algos?

- Implemented in commercial DBMSs
 - DBMSs implement different subsets of known algorithms
- Good algorithms can greatly improve performance
- Need to know about physical operators to understand query optimization

Cost Parameters

- In database systems the data is on disk
- **Cost = total number of I/Os**
- Parameters:
 - $B(R)$ = # of blocks (i.e., pages) for relation R
 - $T(R)$ = # of tuples in relation R
 - $V(R, a)$ = # of distinct values of attribute a
 - When a is a key, $V(R, a) = T(R)$
 - When a is not a key, $V(R, a)$ can be anything $< T(R)$

Cost

- Cost of an operation = number of disk I/Os to
 - Read the operands
 - Compute the result
- Cost of writing the result to disk is *not included*
 - Need to count it separately when applicable

Cost of Scanning a Table

- Result may be unsorted: $B(R)$
- Result needs to be sorted: $3B(R)$
 - We will discuss sorting later

Outline for Today

- **Join operator algorithms**

- One-pass algorithms (Sec. 15.2 and 15.3)
- Index-based algorithms (Sec 15.6)
- Two-pass algorithms (Sec 15.4 and 15.5)
- Note about readings:
 - In class, we will discuss only algorithms for join operator (because other operators are easier)
 - Read the book to get more details about these algos
 - Read the book to learn about algos for other operators

Basic Join Algorithms

- Logical operator:
 - $\text{Product}(\text{pname}, \text{cname}) \bowtie \text{Company}(\text{cname}, \text{city})$
- Propose three physical operators for the join, assuming the tables are in main memory:
 - Hash join
 - Nested loop join
 - Sort-merge join

Hash Join

Hash join: $R \bowtie S$

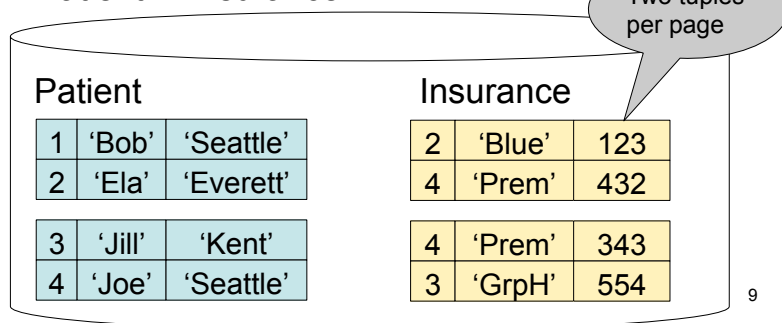
- Scan R, build buckets in main memory
- Then scan S and join
- Cost: $B(R) + B(S)$
- One-pass algorithm when $B(R) \leq M$
 - By “one pass”, we mean that the operator reads its operands only once. It does not write intermediate results back to disk.

Hash Join Example

Patient(pid, name, address)

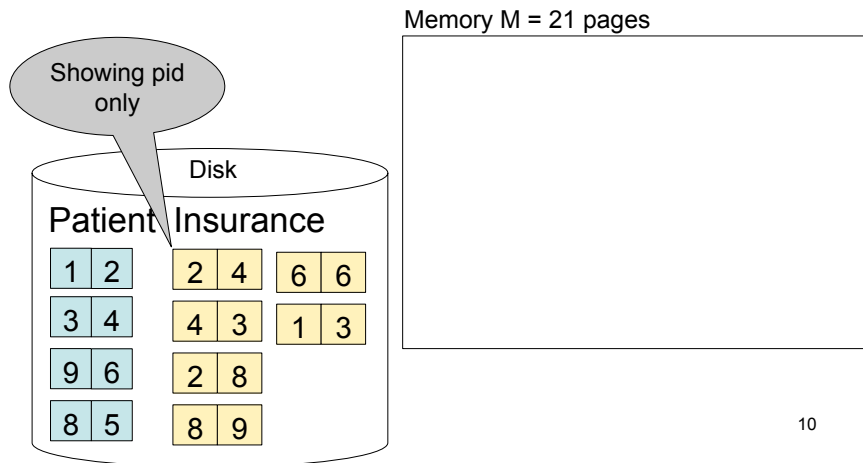
Insurance(pid, provider, policy_nb)

Patient \bowtie Insurance



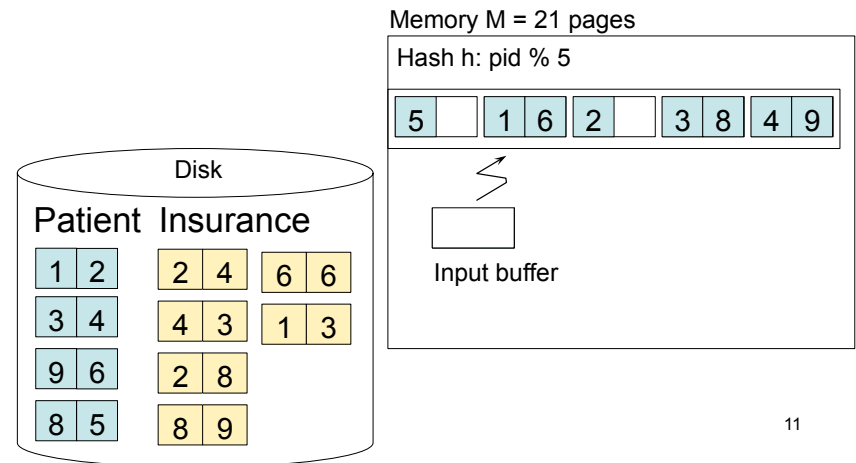
Hash Join Example

Patient \bowtie Insurance



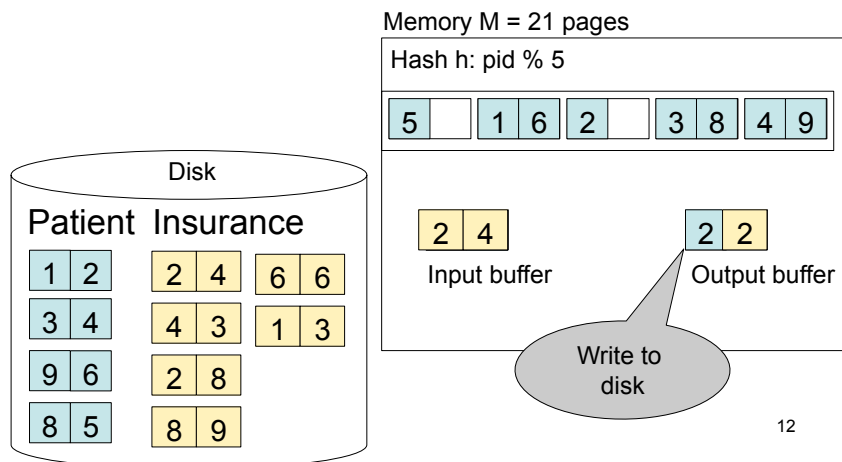
Hash Join Example

Step 1: Scan Patient and create hash table in memory



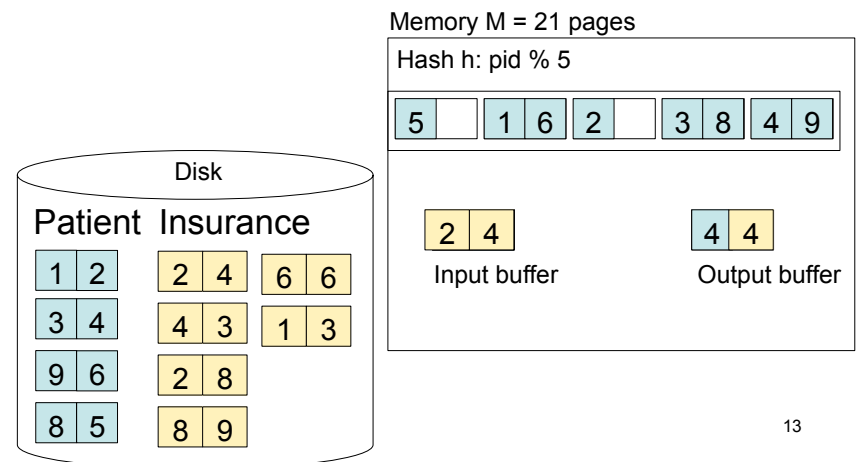
Hash Join Example

Step 2: Scan Insurance and probe into hash table



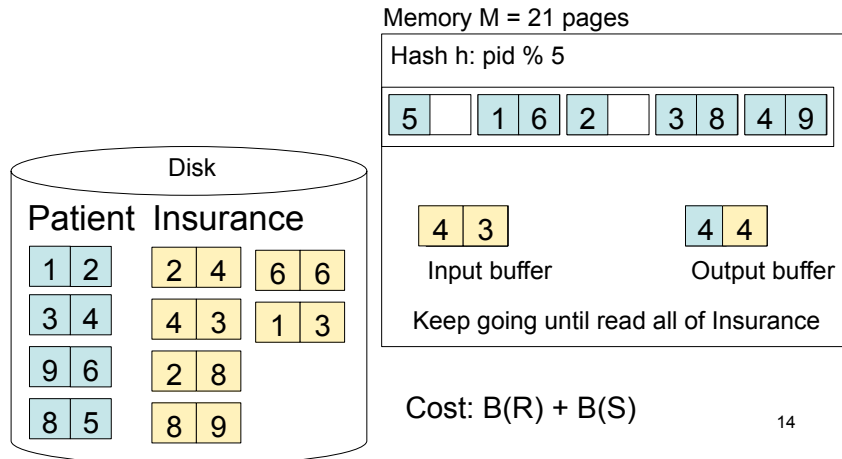
Hash Join Example

Step 2: Scan Insurance and probe into hash table



Hash Join Example

Step 2: Scan Insurance and probe into hash table



Hash Join Details

```
Open( ) {
    H = newHashTable( );
    S.Open( );
    x = S.GetNext( );
    while (x != null) {
        H.insert(x); x = S.GetNext( );
    }
    S.Close( );
    R.Open( );
    buffer = [ ];
}
```

15

Hash Join Details

```
GetNext( ) {
    while (buffer == [ ]) {
        x = R.GetNext( );
        if (x == Null) return NULL;
        buffer = H.find(x);
    }
    z = buffer.first( );
    buffer = buffer.rest( );
    return z;
}
```

16

Hash Join Details

```
Close( ) {
    release memory (H, buffer, etc.);
    R.Close( );
}
```

17

Nested Loop Joins

- Tuple-based nested loop $R \bowtie S$
- R is the outer relation, S is the inner relation

```

for each tuple r in R do
  for each tuple s in S do
    if r and s join then output (r,s)
    
```

- Cost: $B(R) + T(R) B(S)$
- Not quite one-pass since S is read many times

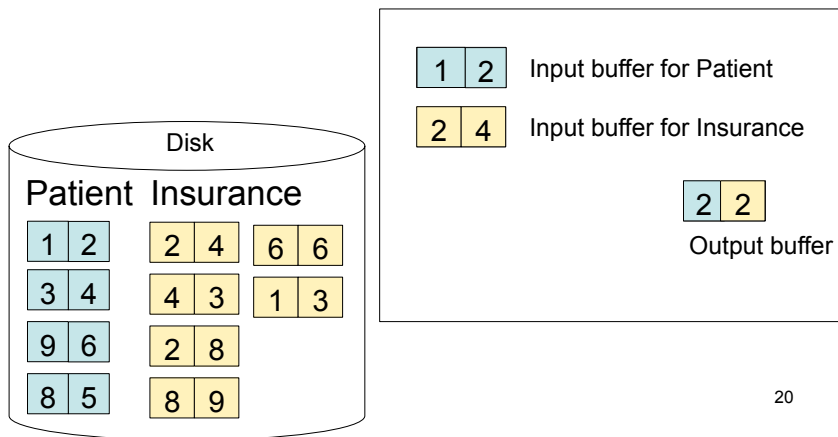
Page-at-a-time Refinement

```

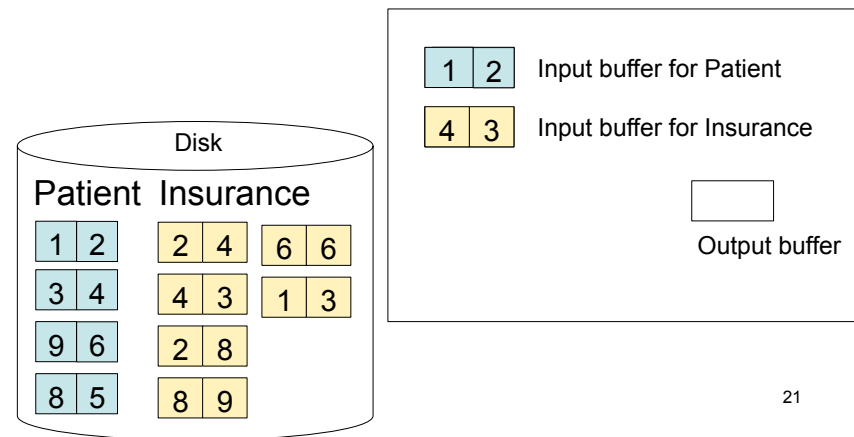
for each page of tuples r in R do
  for each page of tuples s in S do
    for all pairs of tuples
      if r and s join then output (r,s)
    
```

- Cost: $B(R) + B(R)B(S)$

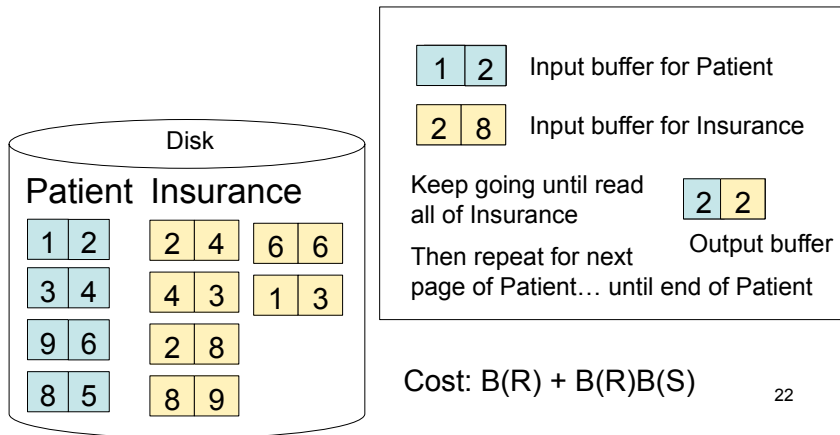
Nested Loop Example



Nested Loop Example



Nested Loop Example



22

Sort-Merge Join

Sort-merge join: $R \bowtie S$

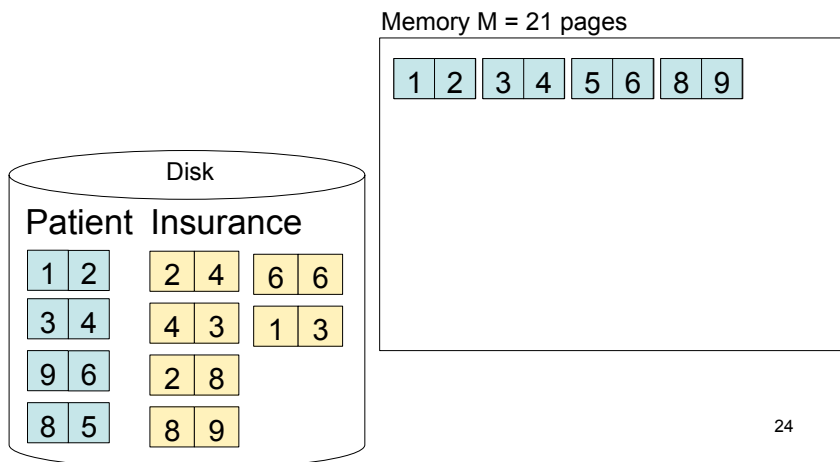
- Scan R and sort in main memory
- Scan S and sort in main memory
- Merge R and S
- Cost: $B(R) + B(S)$
- One pass algorithm when $B(S) + B(R) \leq M$
- Typically, this is NOT a one pass algorithm

Magda Balazinska - CSE 444, Fall 2010

23

Sort-Merge Join Example

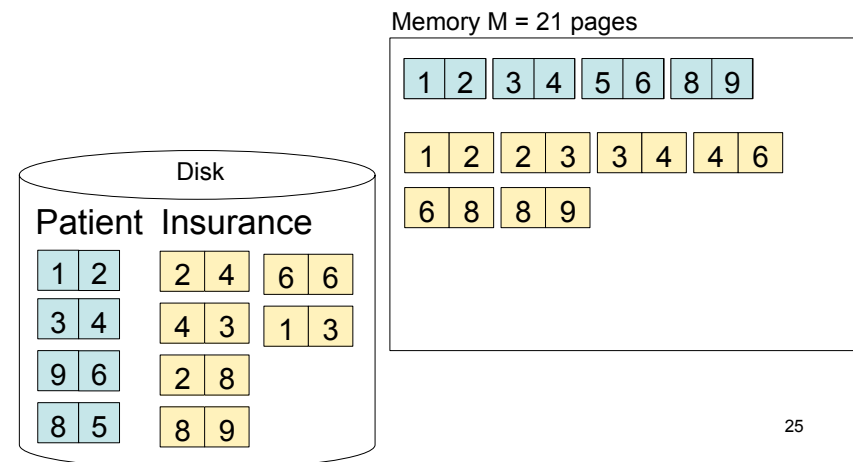
Step 1: Scan Patient and sort in memory



24

Sort-Merge Join Example

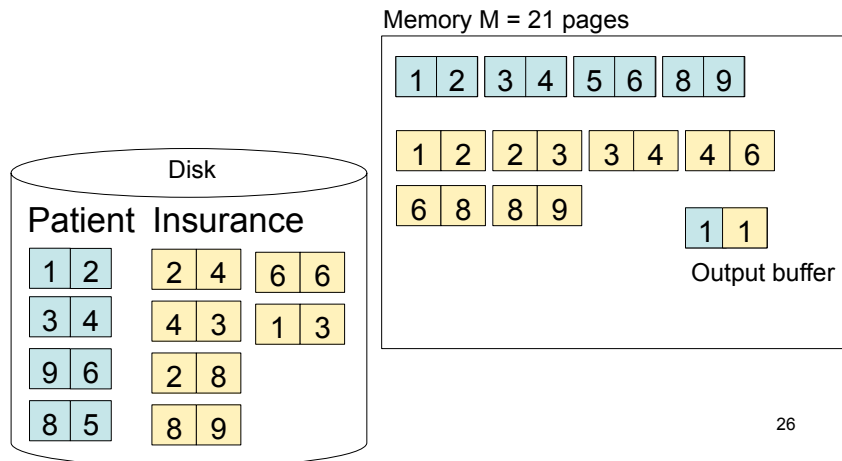
Step 2: Scan Insurance and sort in memory



25

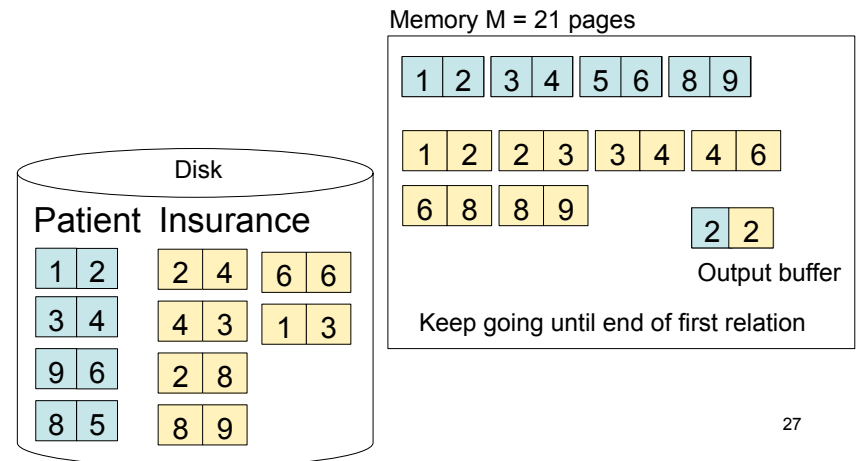
Sort-Merge Join Example

Step 3: Merge Patient and Insurance



Sort-Merge Join Example

Step 3: Merge Patient and Insurance



Outline for Today

- **Join operator algorithms**
 - One-pass algorithms (Sec. 15.2 and 15.3)
 - Index-based algorithms (Sec 15.6)
 - Two-pass algorithms (Sec 15.4 and 15.5)

Review: Access Methods

- **Heap file**
 - Scan tuples one at the time
- **Hash-based index**
 - Efficient selection on equality predicates
 - Can also scan data entries in index
- **Tree-based index**
 - Efficient selection on equality or range predicates
 - Can also scan data entries in index

Index Based Selection

- Selection on equality: $\sigma_{a=v}(R)$
- $V(R, a) = \#$ of distinct values of attribute a
- Clustered index on a : cost $B(R)/V(R,a)$
- Unclustered index on a : cost $T(R)/V(R,a)$
- Note: we ignored I/O cost for index pages

Index Based Selection

- Example:

$B(R) = 2000$ $T(R) = 100,000$ $V(R, a) = 20$

cost of $\sigma_{a=v}(R) = ?$

- Table scan: $B(R) = 2,000$ I/Os
- Index based selection
 - If index is clustered: $B(R)/V(R,a) = 100$ I/Os
 - If index is unclustered: $T(R)/V(R,a) = 5,000$ I/Os
- Lesson
 - Don't build unclustered indexes when $V(R,a)$ is small !

Index Nested Loop Join

$R \bowtie S$

- Assume S has an index on the join attribute
- Iterate over R , for each tuple fetch corresponding tuple(s) from S
- **Cost:**
 - If index on S is clustered: $B(R) + T(R)B(S)/V(S,a)$
 - If index on S is unclustered: $B(R) + T(R)T(S)/V(S,a)$

Outline for Today


- **Join operator algorithms**
 - One-pass algorithms (Sec. 15.2 and 15.3)
 - Index-based algorithms (Sec 15.6)
 - Two-pass algorithms (Sec 15.4 and 15.5)

Two-Pass Algorithms

- What if data does not fit in memory?
- Need to process it in multiple passes
- Two key techniques
 - Hashing
 - Sorting

Two Pass Algorithms Based on Hashing

- Idea: partition a relation R into buckets, on disk
- Each bucket has size approx. $B(R)/M$



- Does each bucket fit in main memory ?
 - Yes if $B(R)/M \leq M$, i.e. $B(R) \leq M^2$


Partitioned (Grace) Hash Join

$R \bowtie S$

- Step 1:
 - Hash S into M-1 buckets
 - Send all buckets to disk
- Step 2
 - Hash R into M-1 buckets
 - Send all buckets to disk
- Step 3
 - Join every pair of buckets

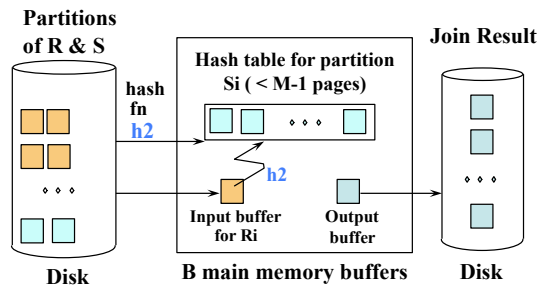
Partitioned Hash Join

- Partition both relations using hash fn **h**
- R tuples in partition i will only match S tuples in partition i.



Partitioned Hash Join

- Read in partition of R, hash it using $h_2 (\neq h)$
 - Build phase
- Scan matching partition of S, search for matches
 - Probe phase



Magda Balazinska - CSE 444, Fall 2010

38

Partitioned Hash Join

- Cost: $3B(R) + 3B(S)$
- Assumption: $\min(B(R), B(S)) \leq M^2$

Magda Balazinska - CSE 444, Fall 2010

39

Partitioned Hash Join

- See detailed example on the board

Magda Balazinska - CSE 444, Fall 2010

40

External Sorting


- Problem: Sort a file of size B with memory M
- Where we need this:
 - ORDER BY in SQL queries
 - Several physical operators
 - Bulk loading of B+-tree indexes.
- Sorting is two-pass when $B < M^2$

Magda Balazinska - CSE 444, Fall 2010

41

External Merge-Sort: Step 1


- Phase one: load M pages in memory, sort



42

External Merge-Sort: Step 2

- Merge $M - 1$ runs into a new run
- Result: runs of length $M(M - 1) \approx M^2$



If $B \leq M^2$ then we are done

Magda Balazinska - CSE 444, Fall 2010

43

External Merge-Sort

- Cost:
 - Read+write+read = $3B(R)$
 - Assumption: $B(R) \leq M^2$
- Other considerations
 - In general, a lot of optimizations are possible

External Merge-Sort

- See detailed example on the board

Two-Pass Join Algorithm Based on Sorting

Join $R \bowtie S$

- Step 1: sort both R and S on the join attribute:
 - Cost: $4B(R)+4B(S)$ (because need to write to disk)
- Step 2: Read both relations in sorted order, match tuples
 - Cost: $B(R)+B(S)$
- Total cost: $5B(R)+5B(S)$
- Assumption: $B(R) \leq M^2$, $B(S) \leq M^2$

Two-Pass Join Algorithm Based on Sorting

Join $R \bowtie S$

- If $B(R) + B(S) \leq M^2$
 - Or if use a priority queue to create runs of length $2|M|$
- If the number of tuples in R matching those in S is small (or vice versa)
- We can compute the join during the merge phase
- Total cost: $3B(R)+3B(S)$

Two-Pass Join Algorithm Based on Sorting

- See detailed example on the board

Summary of Join Algorithms

- **Nested Loop Join**: $B(R) + B(R)B(S)$
 - Assuming page-at-a-time refinement
- **Hash Join**: $3B(R) + 3B(S)$
 - Assuming: $\min(B(R), B(S)) \leq M^2$
- **Sort-Merge Join**: $3B(R)+3B(S)$
 - Assuming $B(R)+B(S) \leq M^2$
- **Index Nested Loop Join**: $B(R) + T(R)B(S)/V(S,a)$
 - Assuming S has clustered index on a