


# B+Tree

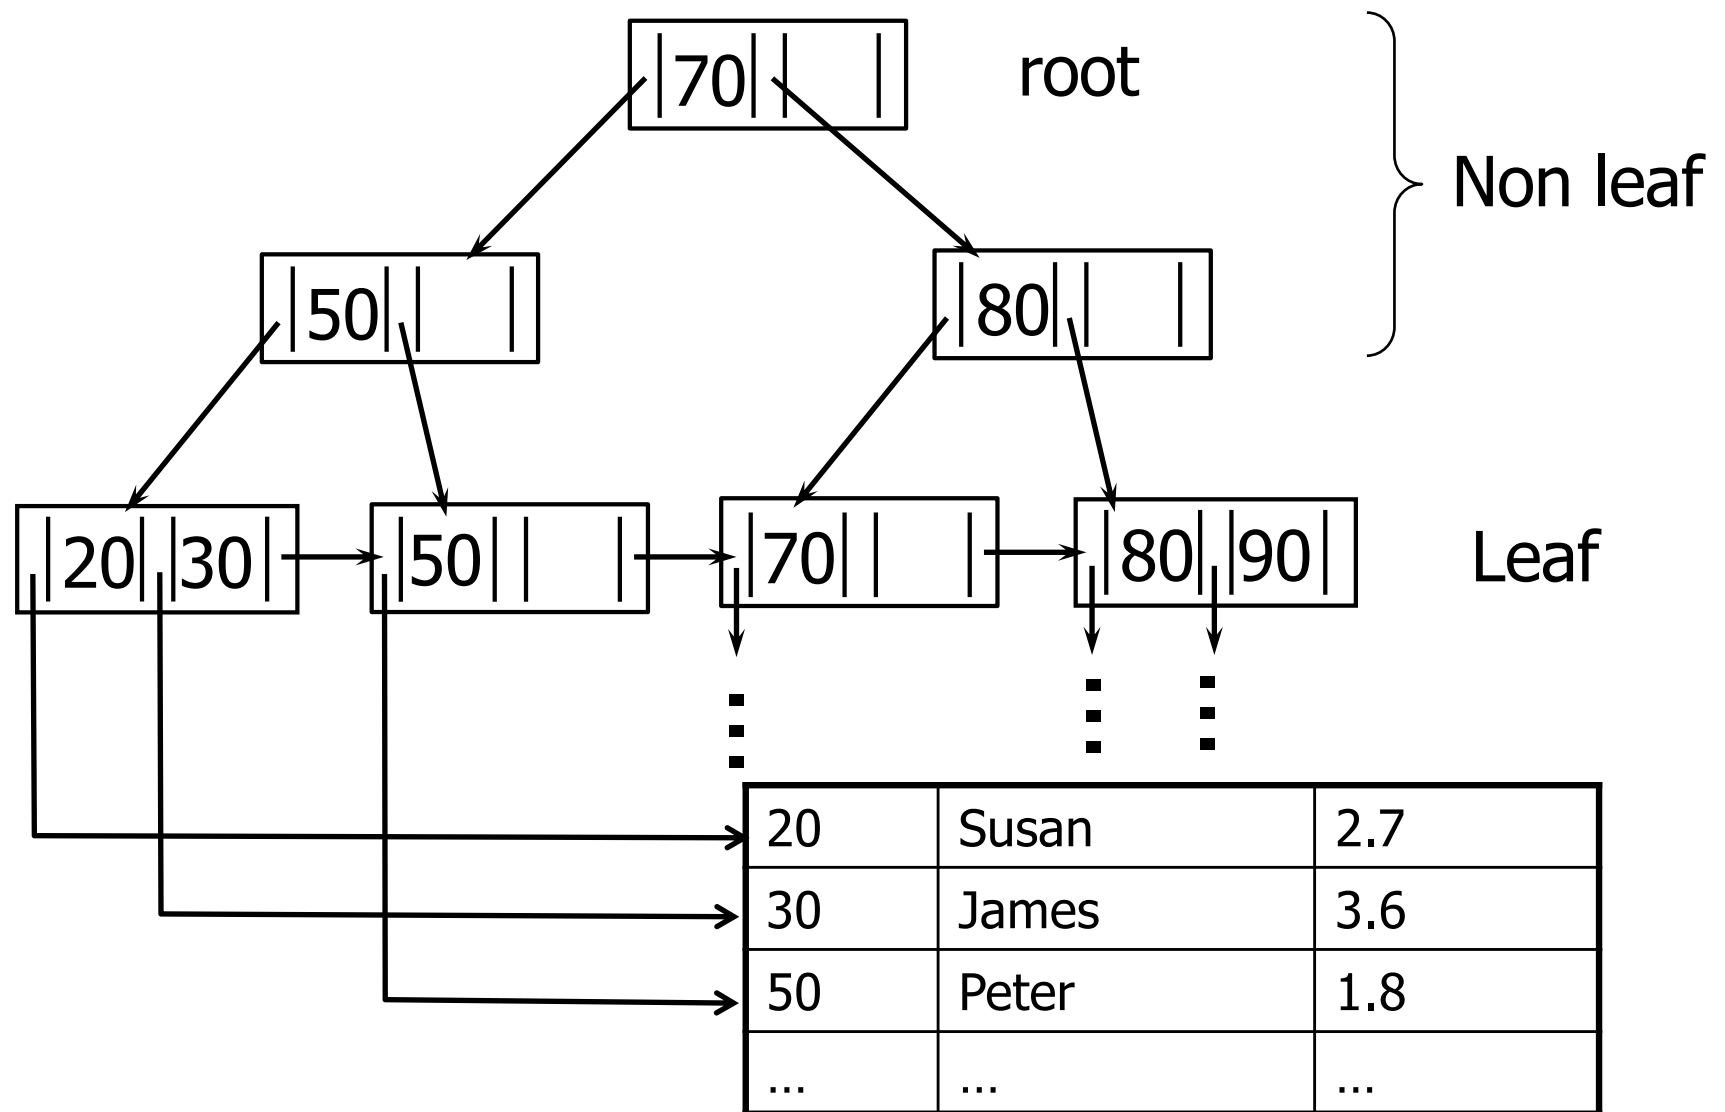
- Most popular index structure in RDBMS
- Advantage
  - Suitable for dynamic updates
  - Balanced
  - Minimum space usage guarantee
- Disadvantage
  - Non-sequential index blocks

B+ tree: generalizing B trees: e.g.,  $n=3$   
in practice much larger



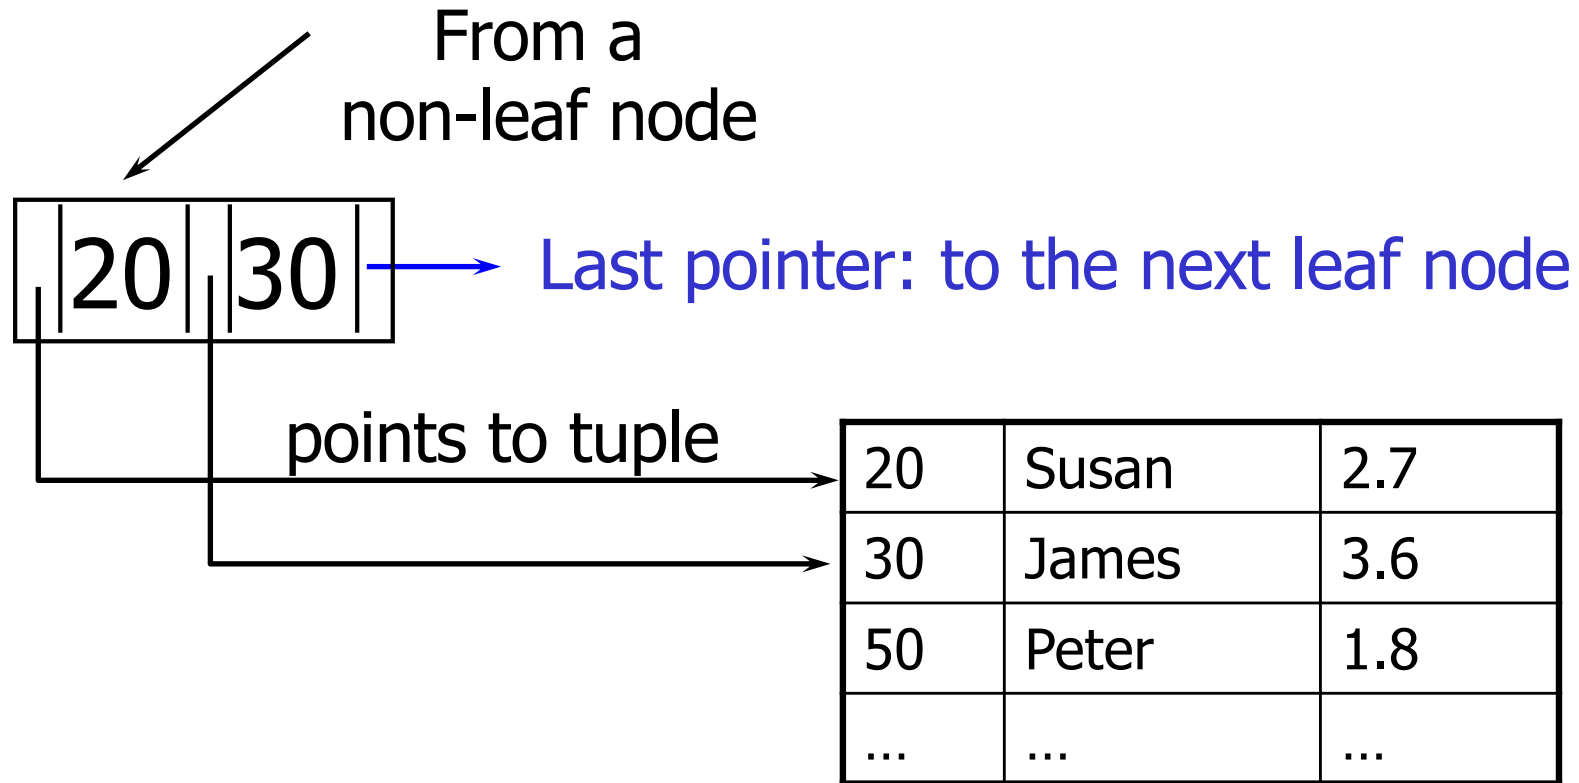
- Points to the nodes one-level below
  - No direct pointers to tuples
- At least half of the ptrs used (precisely,  $\lceil n/2 \rceil$ )
  - except root, where at least 2 ptrs used

# B+Tree Example (n=3)




Balanced: All leaf nodes are at the same level

# Sample Leaf Node (n=3)



- n: max # of pointers in a node
- All pointers (except the last one) point to tuples
- At least half of the pointers are used.  
(more precisely,  $\lceil (n+1)/2 \rceil$  pointers)


### Sample Non-leaf Node (n=3)



- Points to the nodes one-level below
  - No direct pointers to tuples
- At least half of the ptrs used (precisely,  $\lceil n/2 \rceil$ )
  - except root, where at least 2 ptrs used

# Search on B+tree

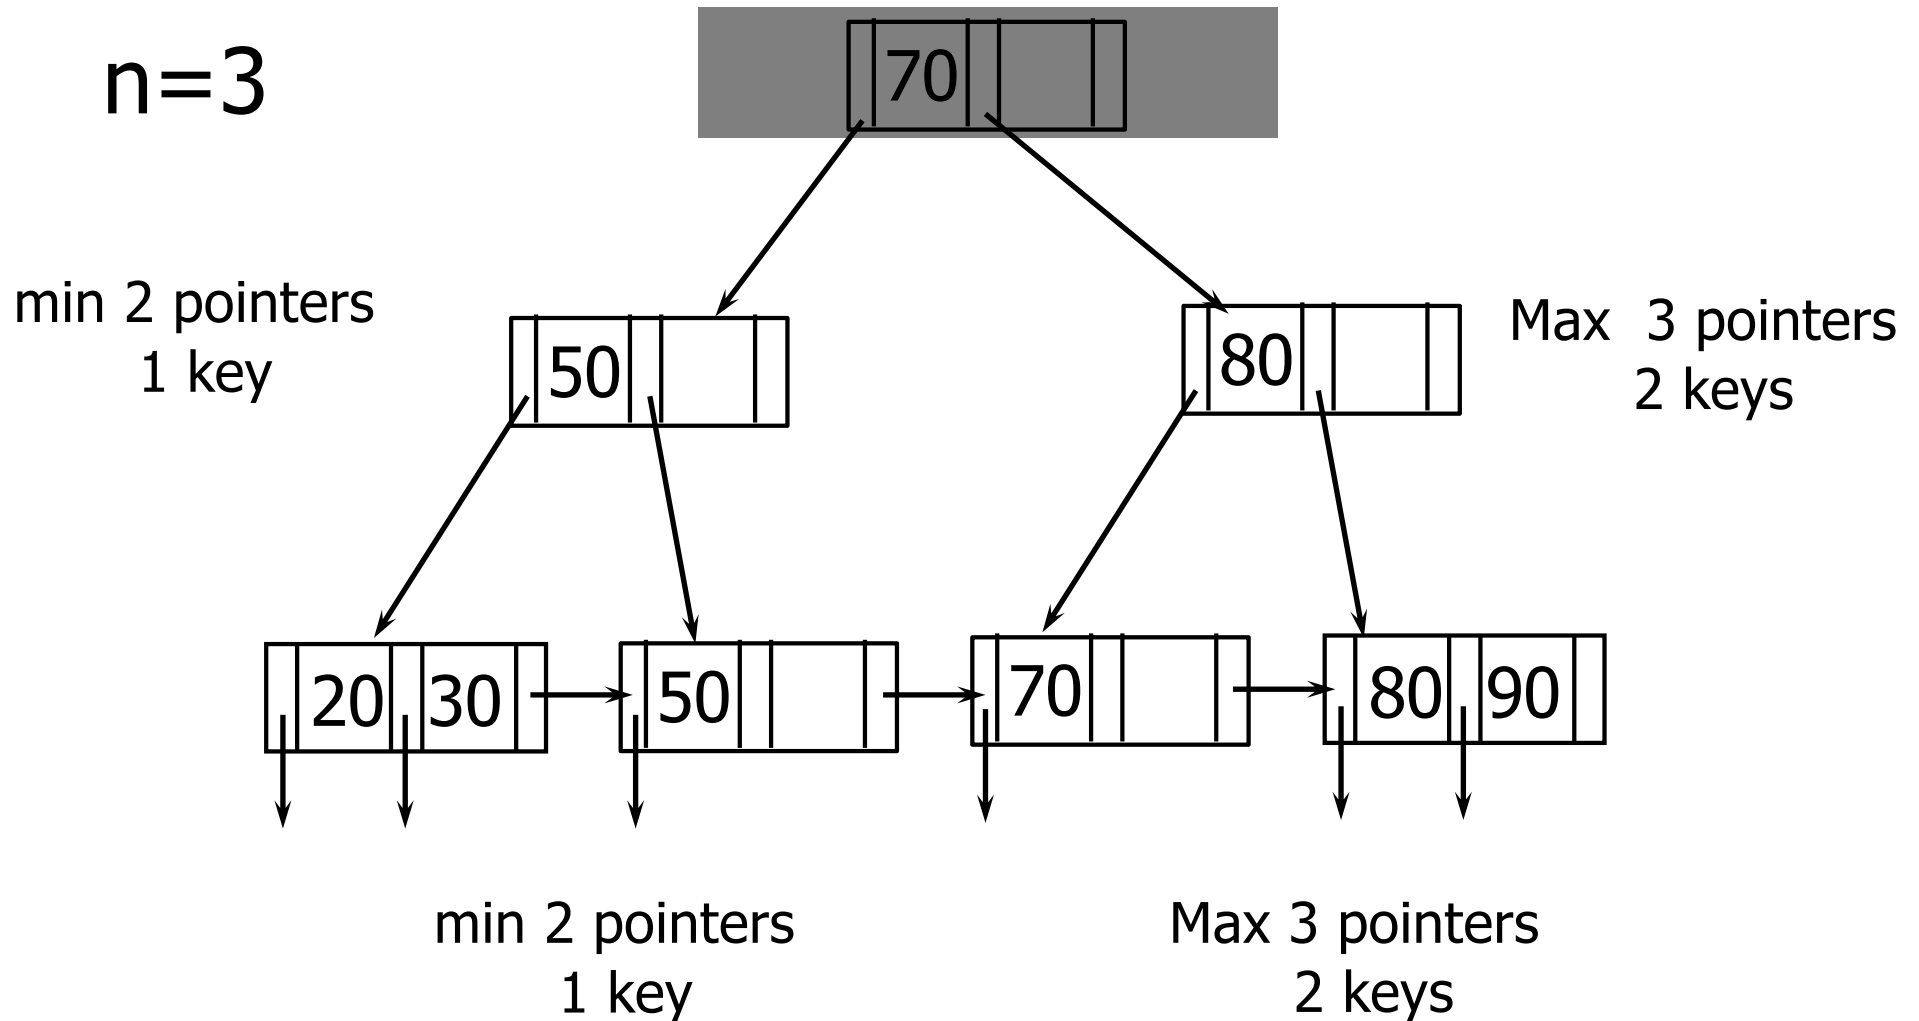
- Find 30, 60, 70?



- Find a greater key and follow the link on the left  
(Algorithm: Figure 12.10 on textbook)

# Nodes but the root are never too empty

$n=3$




## Nodes are never too empty

- Use at least

Non-leaf:  $\lceil n/2 \rceil$  pointers

Leaf:  $\lceil (n+1)/2 \rceil$  pointers





# Number of Ptrs/Keys for B+tree

	Max Ptrs	Max keys	Min ptrs	Min keys
Non-leaf (non-root)	n	n-1	$\lceil n/2 \rceil$	$\lceil n/2 \rceil - 1$
Leaf (non-root)	n	n-1	$\lceil (n+1)/2 \rceil$	$\lceil (n-1)/2 \rceil$
Root	n	n-1	2	1


# B+Tree Insertion

- (a) simple case (no overflow)
- (b) leaf overflow
- (c) non-leaf overflow
- (d) new root

(a) Simple case  
(no overflow)


## Insertion (Simple Case)

- Insert 60



## Insertion (Simple Case)


- Insert 60



(b) Leaf overflow

# Insertion (Leaf Overflow)


- Insert 55



- No space to store 55

# Insertion (Leaf Overflow)

- Insert 55




- Split the leaf into two. Put the keys half and half




# Insertion (Leaf Overflow)

- Insert 55



# Insertion (Leaf Overflow)

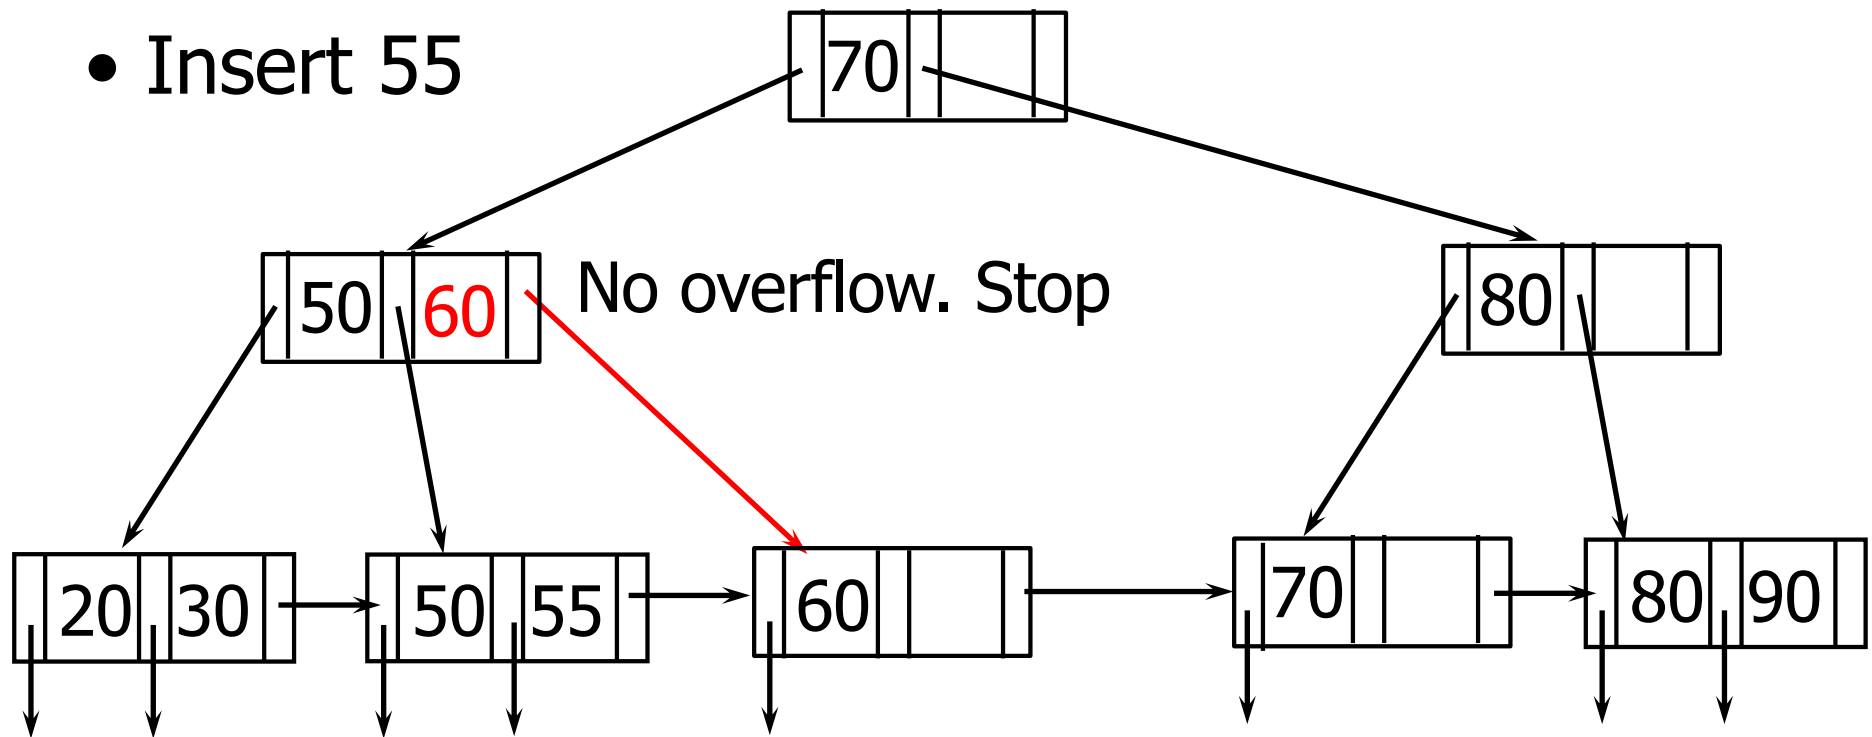
- Insert 55



- Copy the first key of the new node to parent

# Insertion (Leaf Overflow)

- Insert 55




- Q: After split, leaf nodes always half full?

(c) Non-leaf overflow

## Insertion (Non-leaf Overflow)


- Insert 52



Leaf overflow. Split and copy the first key of the new node


# Insertion (Non-leaf Overflow)

- Insert 52



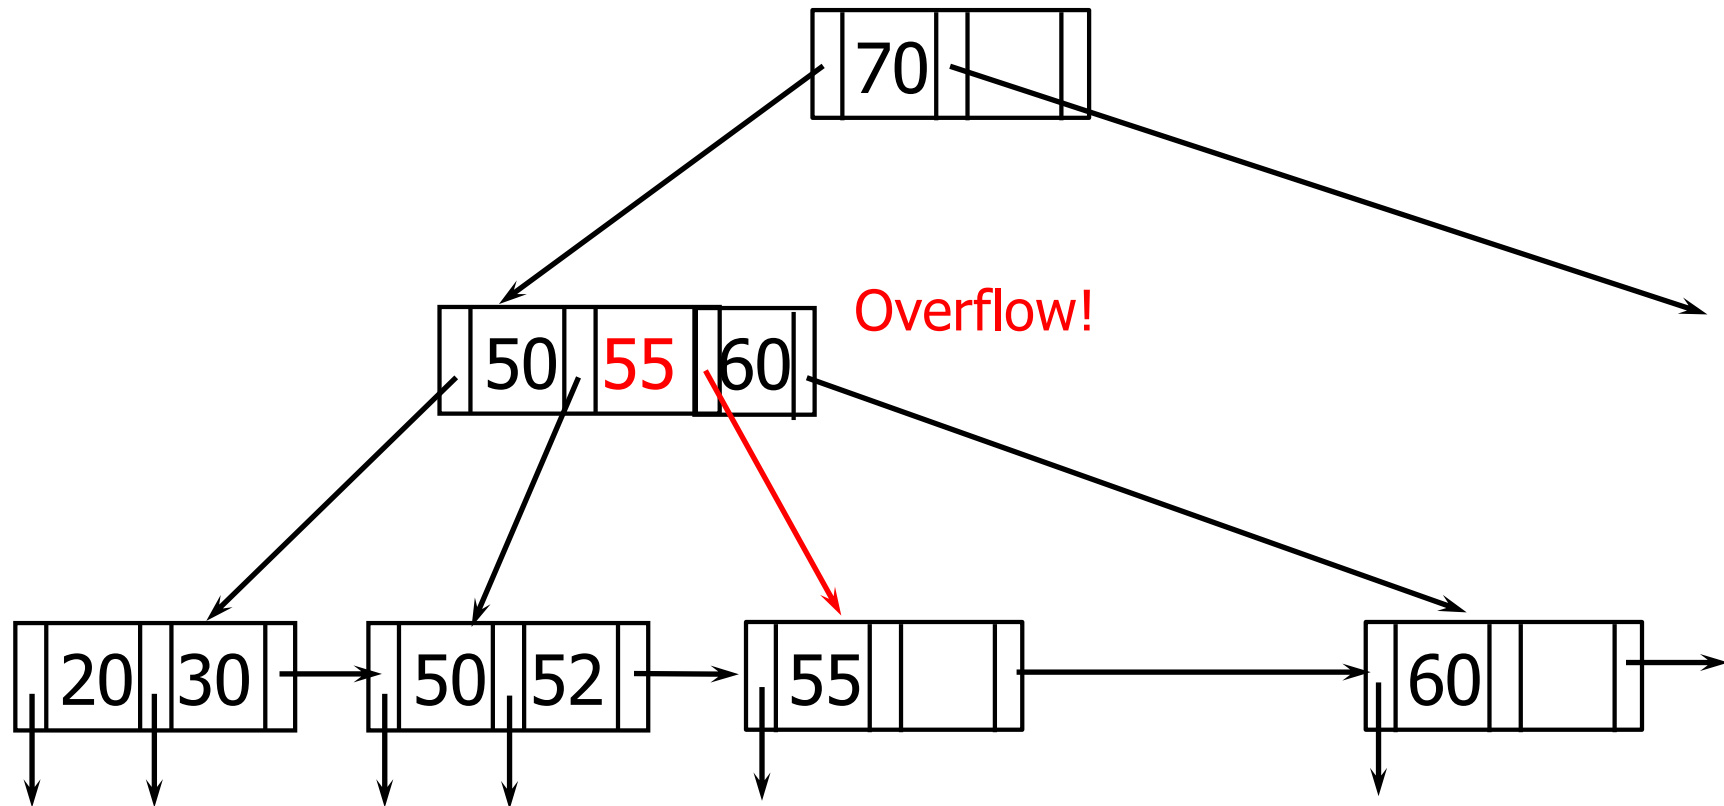
# Insertion (Non-leaf Overflow)

- Insert 52



# Insertion (Non-leaf Overflow)


- Insert 52





# Insertion (Non-leaf Overflow)


- Insert 52



Split the node into two. Move up the key in the middle.


# Insertion (Non-leaf Overflow)

- Insert 52



# Insertion (Non-leaf Overflow)

- Insert 52




Q: After split, non-leaf at least half full?

(d) New root


# Insertion (New Root Node)

- Insert 25



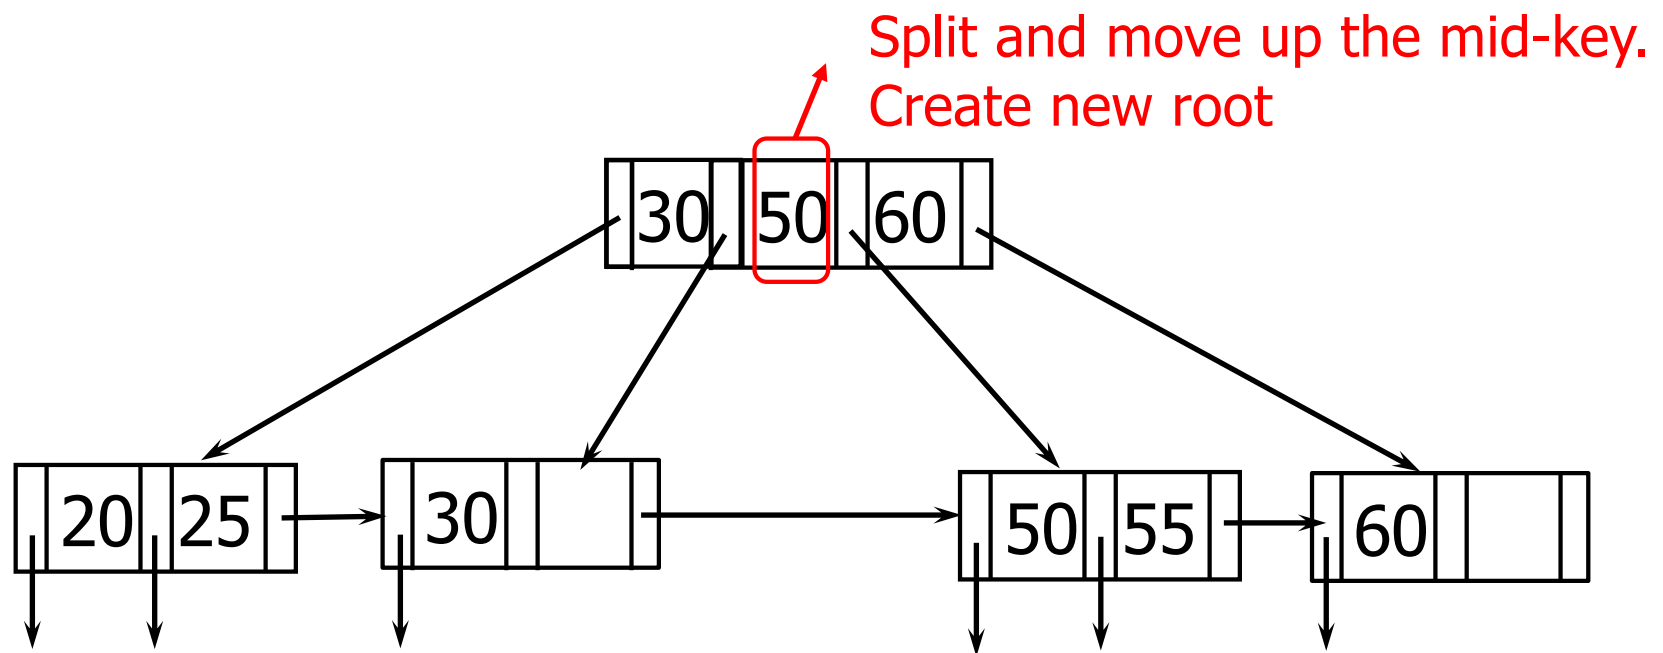
# Insertion (New Root Node)

- Insert 25




# Insertion (New Root Node)

- Insert 25



# Insertion (New Root Node)

- Insert 25
- Q: At least 2 ptrs at root?





# B+Tree Insertion

- Leaf node overflow
  - The first key of the new node is *copied* to the parent
- Non-leaf node overflow
  - The middle key is *moved* to the parent
- Detailed algorithm: see textbook

# B+Tree Deletion


- (a) Simple case (no underflow)
- (b) Leaf node, coalesce with neighbor
- (c) Leaf node, redistribute with neighbor
- (d) Non-leaf node, coalesce with neighbor
- (e) Non-leaf node, redistribute with neighbor

In the examples,  $n = 4$

- Underflow for non-leaf when fewer than  $\lceil n/2 \rceil = 2$  ptrs
- Underflow for leaf when fewer than  $\lceil (n+1)/2 \rceil = 3$  ptrs
- Nodes are labeled as  $a, b, c, d, \dots$


(a) Simple case  
(no underflow)

## (a) Simple case



- Delete 25


## (a) Simple case



- Delete 25
  - Underflow? Min 3 ptrs. Currently 3 ptrs


(b) Leaf node, coalesce  
with neighbor

## (b) Coalesce with sibling (leaf)



- Delete 50


## (b) Coalesce with sibling (leaf)



- Delete 50
  - Underflow? Min 3 ptrs, currently 2.




## (b) Coalesce with sibling (leaf)




- Delete 50
  - Try to merge with a sibling

## (b) Coalesce with sibling (leaf)




- Delete 50
  - Merge *c* and *d*. Move everything on the right to the left.

## (b) Coalesce with sibling (leaf)




- Delete 50
  - Once everything is moved, delete *d*

## (b) Coalesce with sibling (leaf)



- Delete 50
  - After leaf node merge,
    - From its parent, delete the pointer and key to the deleted node


## (b) Coalesce with sibling (leaf)



- Delete 50
  - Check underflow at *a*. Min 2 ptrs, currently 3


(c) Leaf node, redistribute  
with neighbor

## (c) Redistribute (leaf)



- Delete 50


## (c) Redistribute (leaf)



- Delete 50
  - Underflow? Min 3 ptrs, currently 2
  - Check if *d* can be merged with its sibling *c* or *e*
  - If not, redistribute the keys in *d* with a sibling
    - Say, with *c*




## (c) Redistribute (leaf)




- Delete 50
  - Redistribute *c* and *d*, so that nodes *c* and *d* are roughly “half full”
    - Move the key 30 and its tuple pointer to the *d*

## (c) Redistribute (leaf)



- Delete 50
  - Update the key in the parent


### (c) Redistribute (leaf)



- Delete 50
  - No underflow at *a*. Done.


(d) Non-leaf node, coalesce  
with neighbor

## (d) Coalesce (non-leaf)




- Delete 20
  - Underflow! Merge *d* with *e*.
    - Move everything in the right to the left

## (d) Coalesce (non-leaf)




- Delete 20
  - From the parent node, delete pointer and key to the deleted node

## (d) Coalesce (non-leaf)



- Delete 20
  - Underflow at *b*? Min 2 ptrs, currently 1.
  - Try to merge with its sibling.
    - Nodes *b* and *c*: 3 ptrs in total. Max 4 ptrs.
    - Merge *b* and *c*.


## (d) Coalesce (non-leaf)



- Delete 20
  - Merge *b* and *c*
    - Pull down the mid-key 50 in the parent node
    - Move everything in the right node to the left.
- Very important: when we merge non-leaf nodes, we always pull down the mid-key in the parent and place it in the merged node.




## (d) Coalesce (non-leaf)




- Delete 20
  - Merge *b* and *c*
    - Pull down the mid-key 50 in the parent node
    - Move everything in the right node to the left.
- Very important: when we merge non-leaf nodes, we always pull down the mid-key in the parent and place it in the merged node.

## (d) Coalesce (non-leaf)



- Delete 20
  - Delete pointer to the merged node.


## (d) Coalesce (non-leaf)



- Delete 20
  - Underflow at *a*? Min 2 ptrs. Currently 2. Done.


(e) Non-leaf node, redistribute  
with neighbor

## (e) Redistribute (non-leaf)




- Delete 20
  - Underflow! Merge *d* with *e*.

## (e) Redistribute (non-leaf)




- Delete 20
  - After merge, remove the key and ptr to the deleted node from the parent

## (e) Redistribute (non-leaf)



- Delete 20
  - Underflow at *b*? Min 2 ptrs, currently 1.
  - Merge *b* with *c*? Max 4 ptrs, 5 ptrs in total.
  - If cannot be merged, redistribute the keys with a sibling.
    - Redistribute *b* and *c*

## (e) Redistribute (non-leaf)




- Delete 20

Redistribution at a non-leaf node is done in two steps.

*Step 1:* Temporarily, make the left node *b* "overflow" by pulling down the mid-key and moving everything to the left.



## (e) Redistribute (non-leaf)




- Delete 20

*Step 2:* Apply the "overflow handling algorithm" (the same algorithm used for B+tree insertion) to the overflowed node

- Detailed algorithm in the next slide

## (e) Redistribute (non-leaf)




- Delete 20

*Step 2: "overflow handling algorithm"*

- Pick the mid-key (say 90) in the node and move it to parent.
- Move everything to the right of 90 to the empty node *c*.

## (e) Redistribute (non-leaf)



- Delete 20
  - Underflow at *a*? Min 2 ptrs, currently 3. Done

# Important Points

- Remember:
  - For leaf node merging, we delete the mid-key from the parent
  - For non-leaf node merging/redistribution, we pull down the mid-key from their parent.
- In practice
  - Coalescing is often not implemented
    - Too hard and not worth it

# Where does $n$ come from?

- $n$  determined by
  - Size of a node
  - Size of search key
  - Size of an index pointer
- Q: 1024B node, 10B key, 8B ptr  $\rightarrow n$ ?
- Computation: 8 bytes for final pointer. Then  $1024 - 8 = 1016$  bytes left for Key+pointer pair, each taking 18 bytes. Thus  $\text{floor}(1016/18) = 56$ . Thus  $n = 57$ .
- Or use formulas

## Summary on tree index

- Issues to consider
  - Sparse vs. dense
  - Primary (clustering) vs. secondary (non-clustering)
- Indexed sequential file (ISAM)
  - Simple algorithm. Sequential blocks
  - Not suitable for dynamic environment
- B+trees
  - Balanced, minimum space guarantee
  - Insertion, deletion algorithms

# Index Creation in SQL

- `CREATE INDEX <indexname>  
ON <table> (<attr>, <attr>, ...)`
- **Example**
  - `CREATE INDEX stidx ON  
Student(sid)`
    - Creates a B+tree on the attributes
    - Speeds up lookup on sid

# Primary (Clustering) Index

- MySQL:
  - Primary key becomes the clustering index
- DB2:
  - `CREATE INDEX idx ON Student(sid)`  
`CLUSTER`
  - Tuples in the table are sequenced by sid
- Oracle: Index-Organized Table (IOT)
  - `CREATE TABLE T (`  
`...`  
`) ORGANIZATION INDEX`
  - B+tree on primary key
  - Tuples are stored at the leaf nodes of B+tree
- Periodic reorganization may still be necessary to improve range scan performance



# Next topic

- Hash index
  - Static hashing
  - Extendible hashing

# What is a Hash Table?

- Hash Table
  - Hash function
    - $h(k)$ : key  $\rightarrow$  integer  $[0\dots n]$
    - e.g.,  $h(\text{'Susan'}) = 7$
  - Array for keys:  $T[0\dots n]$
  - Given a key  $k$ , store it in  $T[h(k)]$

$h(\text{Susan}) = 4$   
 $h(\text{James}) = 3$   
 $h(\text{Neil}) = 1$

0	
1	Neil
2	
3	James
4	Susan
5	

# Overflow and Chaining

- Insert

$h(a) = 1$

$h(b) = 2$

$h(c) = 1$


$h(d) = 0$

$h(e) = 1$

- Delete


$h(b) = 2$

$h(c) = 1$




# Major Problem of Static Hashing

- How to cope with growth?
  - Data tends to grow in size
  - Overflow blocks unavoidable




# Extendible Hashing (two ideas)

(a) Use  $i$  of  $b$  bits output by hash function



# Extendible Hashing (two ideas)


(b) Use directory that maintains pointers to hash buckets (indirection)



# Example


- $h(k)$  is 4 bits; 2 keys/bucket

Insert 0111



# Example

Insert 1010

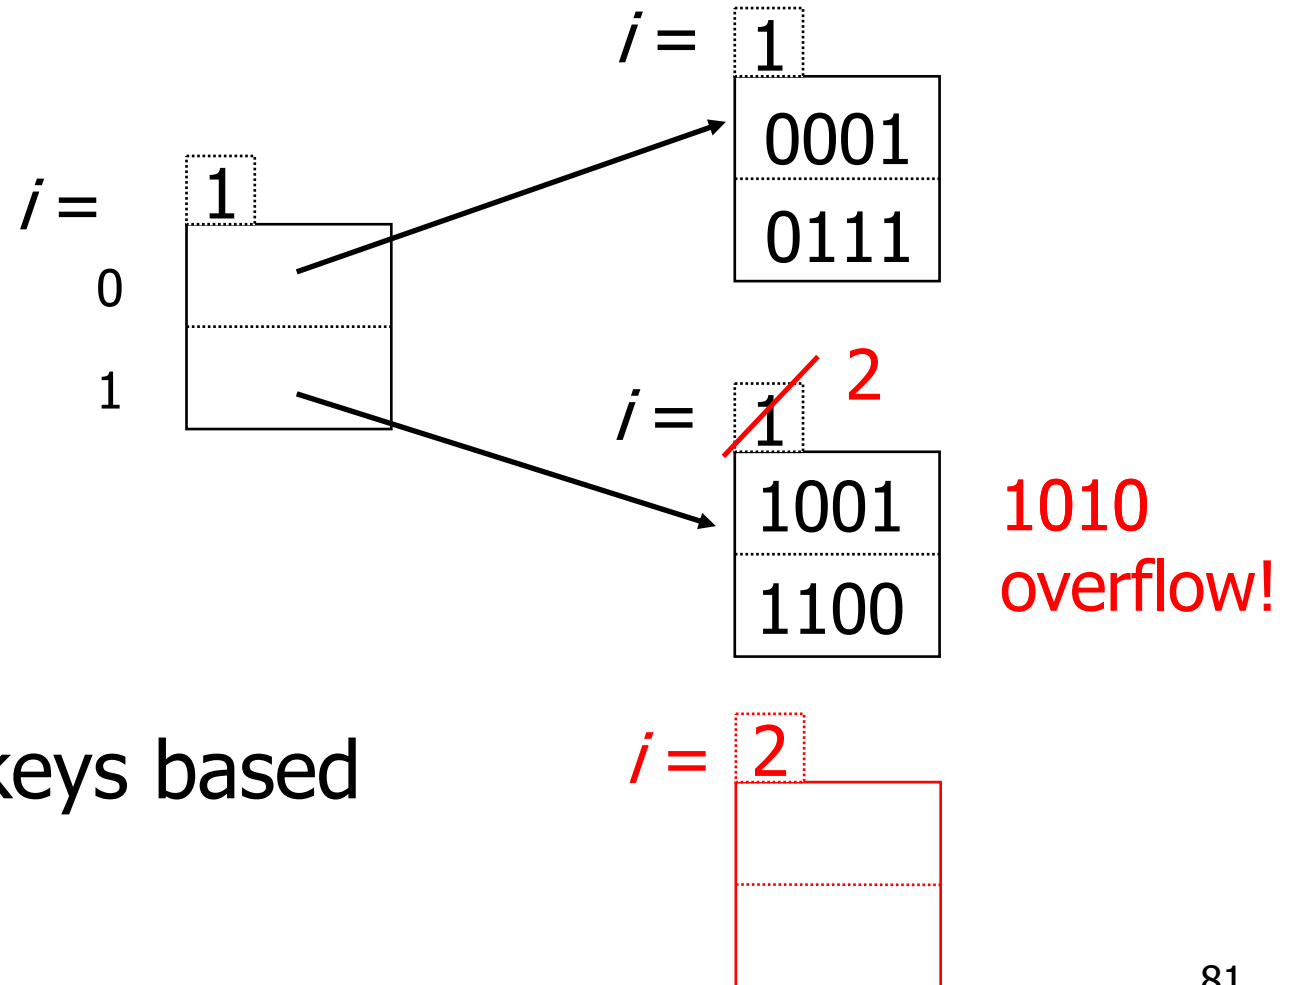


Increase  $i$  of the bucket. Split it.




# Example

Insert 1010



# Example

Insert 1010




Update ptr in dir to new bkt


If no space,  
double directory size (increase  $i$ )

# Example

Insert 1010



Copy pointers



# Example

Insert 1010

$i = 2$

00

01

10

11

$i =$

0

1

$i = 1$

$i = 1$

0001

0111

$i = 2$

1001

1010

$i = 2$

1100

# Example

Insert 0000



$i =$  2

00



01

10

11






0000  
Overflow!





Split bucket and increase  $i$

# Example

Insert 0000








0000  
Overflow!



Redistribute keys

# Example

Insert 0000



Update ptr in directory

# Example

Insert 0000





i = 2

00

01

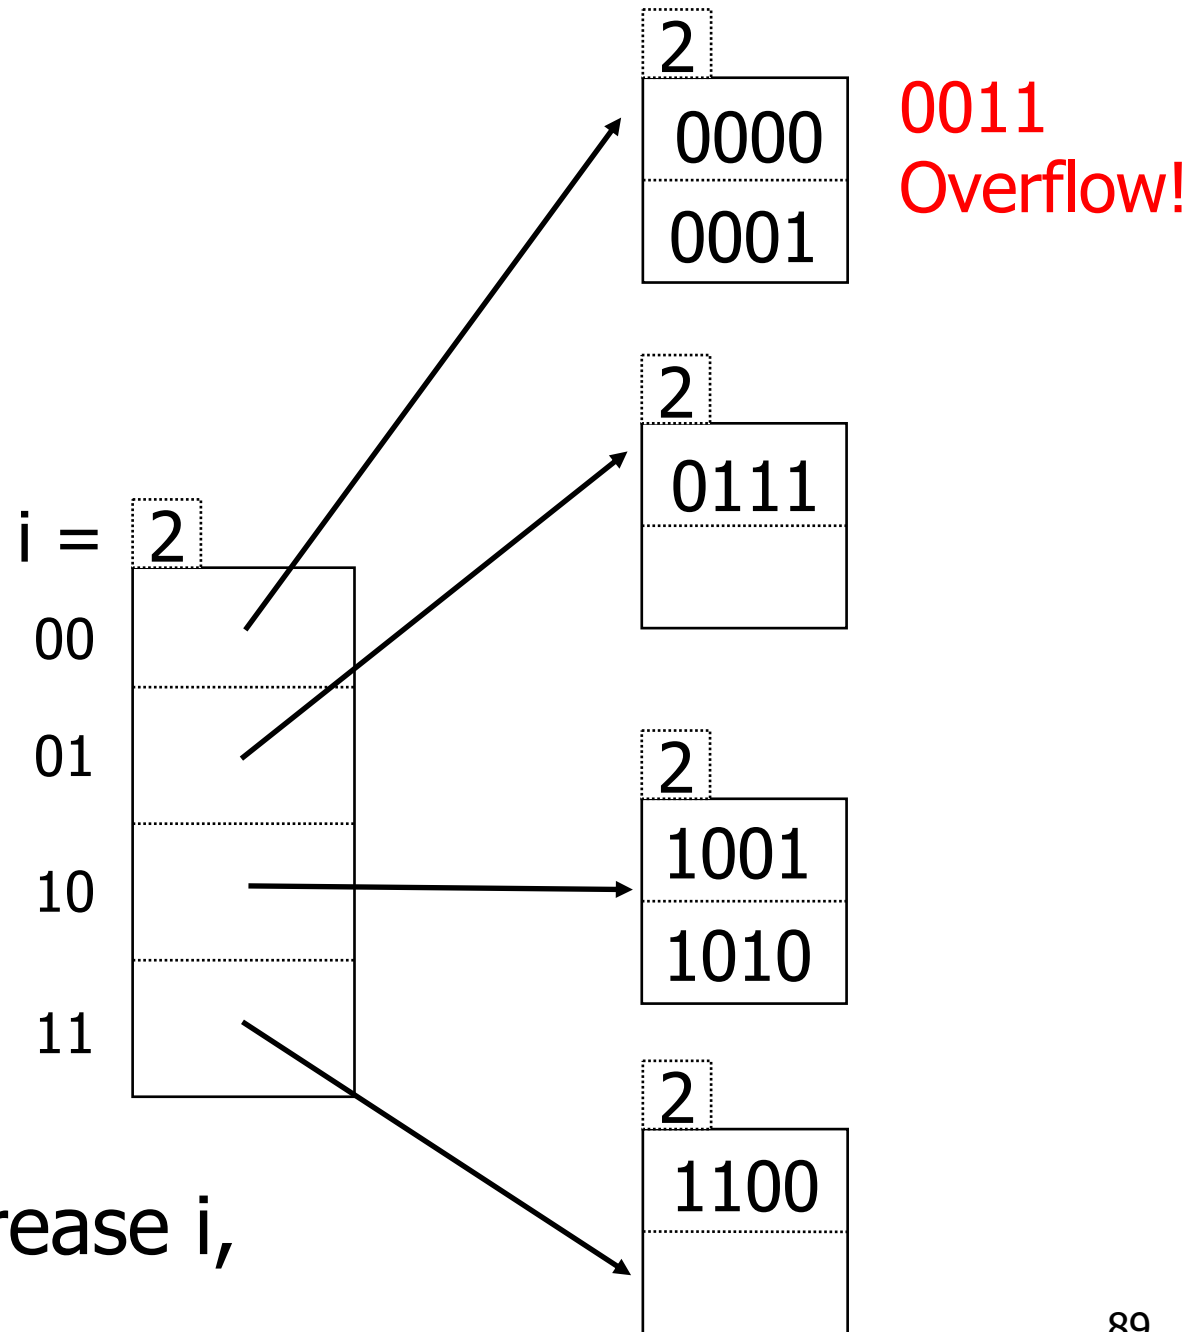
10

11






Insert 0011



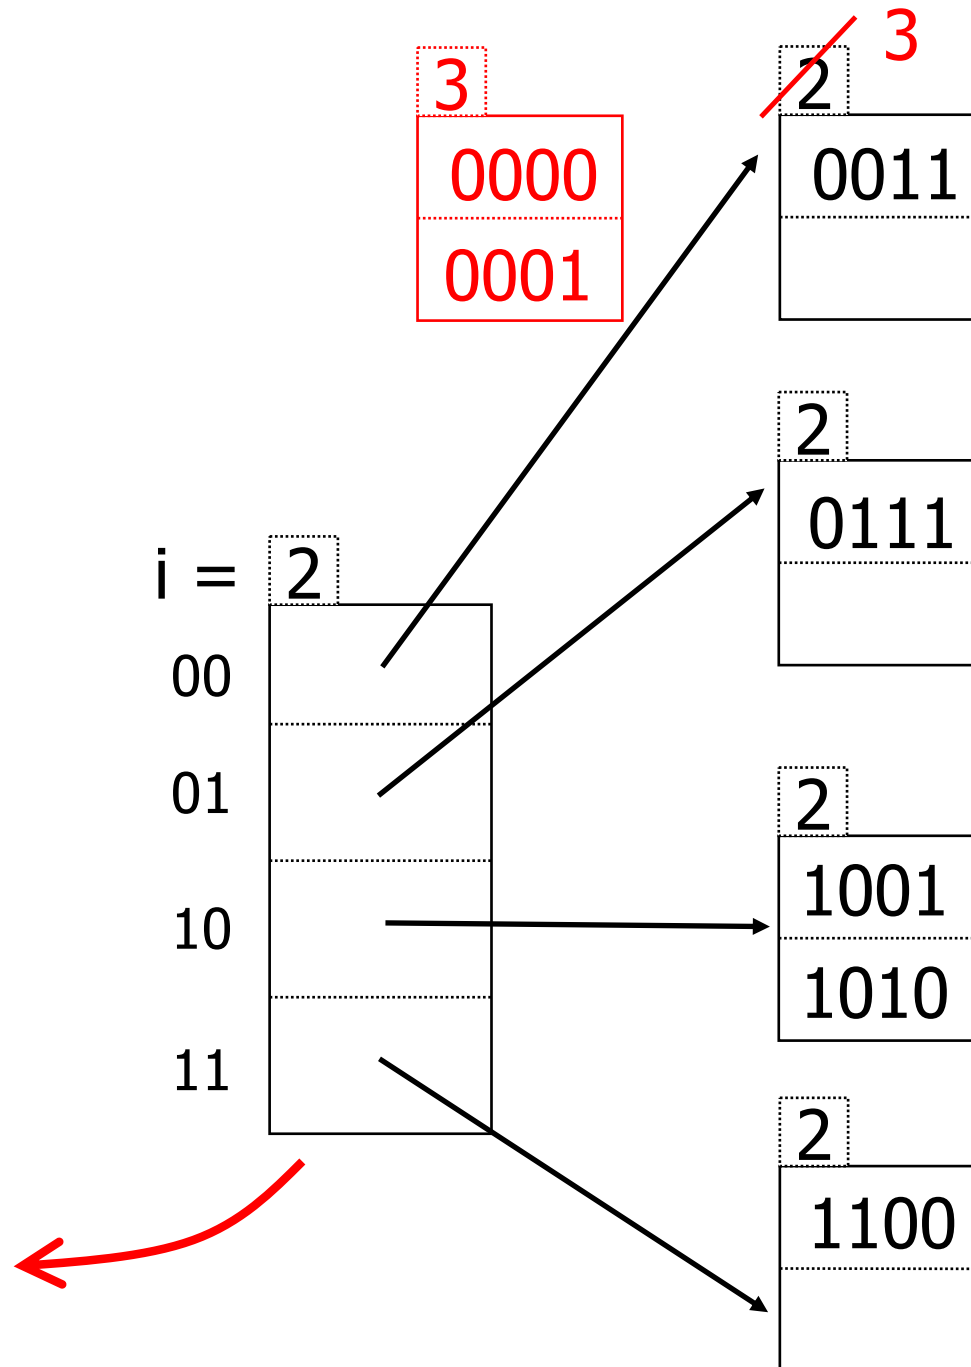

Split bucket, increase  $i$ ,  
redistribute keys

Insert 0011

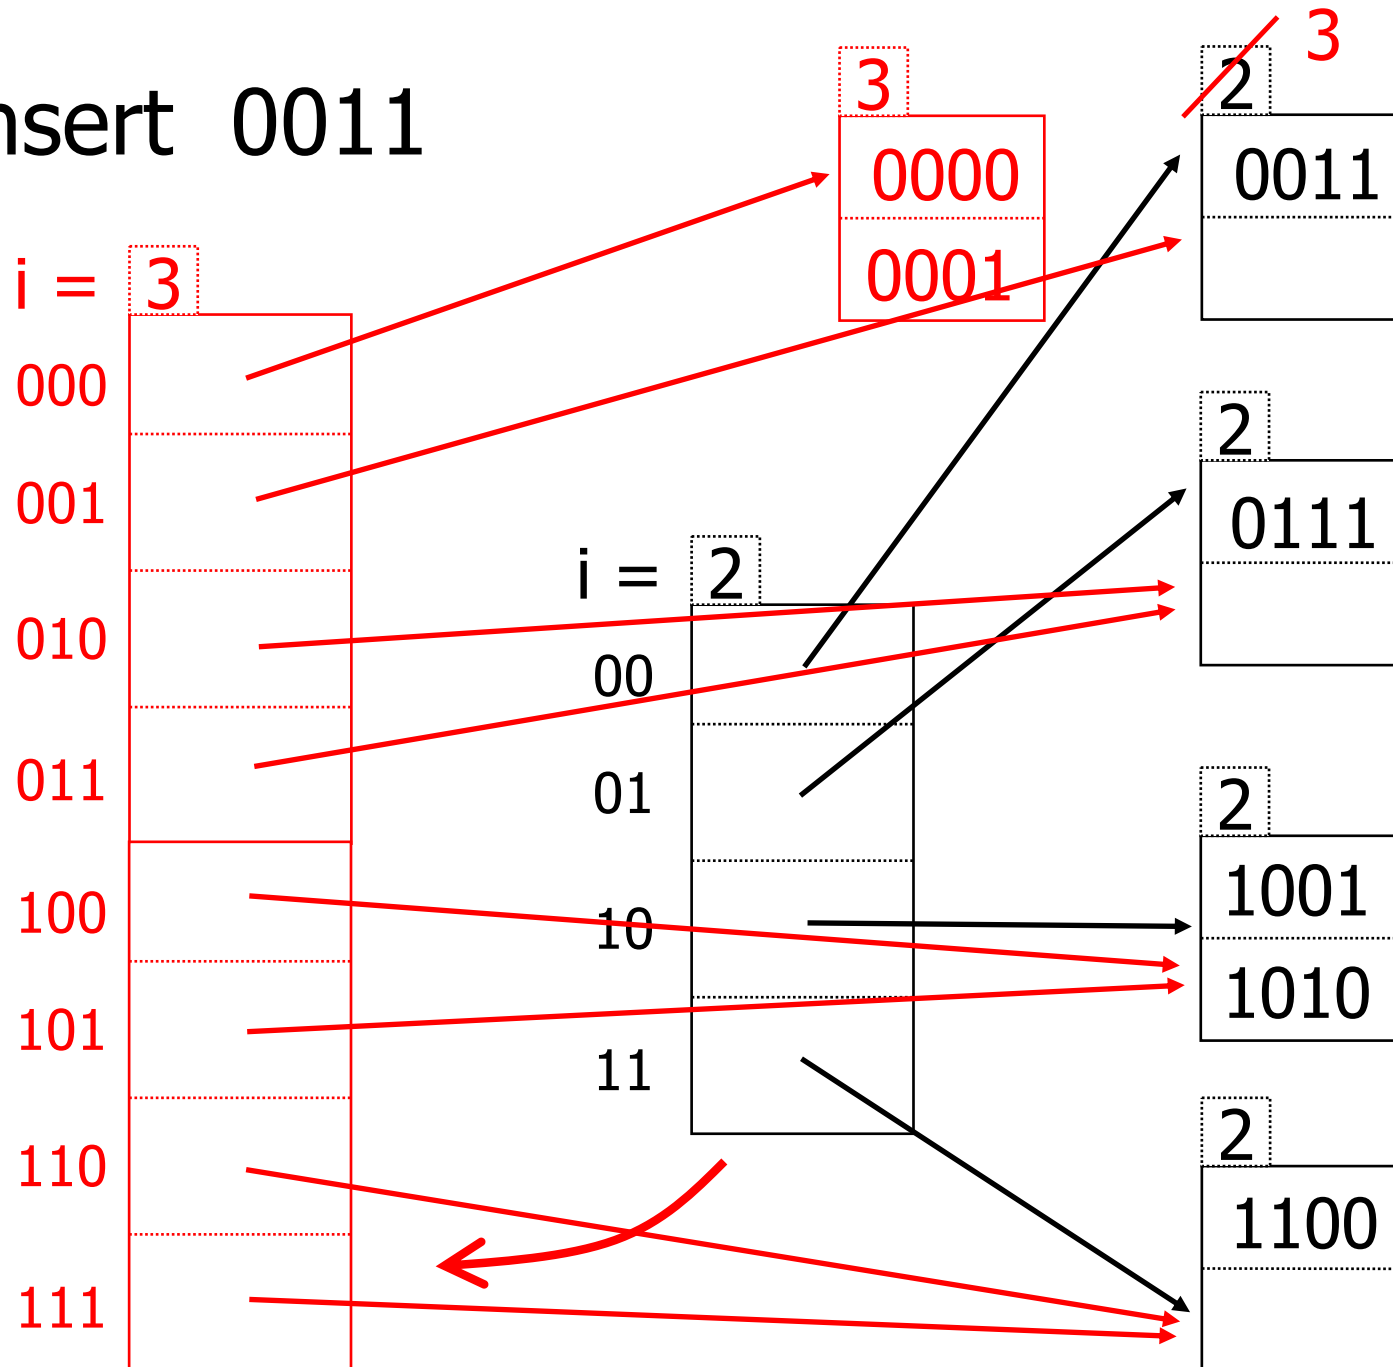


Update ptr in dir  
If no space, double directory

# Insert 0011




# Insert 0011




# Extendible Hashing: Deletion

- Two options
  - a) No merging of buckets
  - b) Merge buckets and shrink directory if possible

# Delete 1010




# Delete 1010



- Can we merge a and b? b and c?

## Delete 1010




Decrease i and merge buckets

Update ptr in directory

Q: Can we shrink directory?



# Delete 1010




# Bucket Merge Condition

- Bucket merge condition
  - Bucket  $i$ 's are the same
  - First  $(i-1)$  bits of the hash key are the same
- Directory shrink condition
  - All bucket  $i$ 's are smaller than the directory  $i$

# Questions on Extendible Hashing

- Can we provide minimum space guarantee?

# Space Waste




# Hash index summary

- Static hashing
  - Overflow and chaining
- Extendible hashing
  - Can handle growing files
    - No periodic reorganizations
  - Indirection
    - Up to 2 disk accesses to access a key
  - Directory doubles in size
    - Not too bad if the data is not too large

## Question on B+tree

- **SELECT \***  
**FROM Student**  
**WHERE sid > 60?**



# Hashing vs. Tree

- Can an extendible-hash index support?

```
SELECT *  
FROM R  
WHERE R.A > 5
```

- Which one is better, B+tree or Extendible hashing?

```
SELECT *  
FROM R  
WHERE R.A = 5
```