

# R JOIN S?

R

<i>c</i>	
40	T1
60	T2
30	T3
10	T4
20	T5

S

<i>c</i>	
10	T6
60	T7
40	T8
20	T9

# Nested-Loop Join (NLJ)

For each  $r \in R$  do

For each  $s \in S$  do

if  $r.C = s.C$  then output  $r,s$  pair

R

40	T1
60	T2
30	T3
10	T4
20	T5

S

10	T6
60	T7
40	T8
20	T9

- $R$  is called the **outer relation** and  $S$  the **inner relation** of the join.
- Requires no indices and can be used with any kind of join condition.
- Expensive since it examines every pair of tuples in the two relations.

## Nested-Loop Join (Cont.)

For each  $r \in R$  do  
  For each  $s \in S$  do  
    if  $r.C = s.C$  then output  $r,s$  pair

- In the worst case, if there is enough memory only to hold one block of each relation, the estimated cost is

$$b_R + n_R * b_S \text{ block transfers}$$

- If the smaller relation fits entirely in memory, use that as the inner relation.
  - Reduces cost to  $b_r + b_s$  block transfers and 2 seeks
- When neither relation fits in memory, **Block Nested-Loops** algorithm (next slide) should be used instead.

# Block Nested-Loop Join

- Variant of nested-loop join in which every block of inner relation is paired with every block of outer relation.

**for each** block  $B_r$  **of**  $r$  **do begin**

**for each** block  $B_s$  **of**  $s$  **do begin**

**for each** tuple  $t_r$  **in**  $B_r$  **do begin**

**for each** tuple  $t_s$  **in**  $B_s$  **do begin**

                Check if  $(t_r, t_s)$  satisfy the join condition

                if they do, add  $t_r \cdot t_s$  to the result.

**end**

**end**

**end**

**end**

# Block Nested-Loop Join (Cont.)


- **Worst case estimate:  $b_R + b_R * b_S$  block transfers**
  - Each block in the inner relation  $S$  is read once for each *block* in the outer relation  $R$ .
- **Best case:  $b_R + b_S$  block transfers** (*if  $1 + b_S$  blocks fit in memory*)
- Many Improvements proposed to nested loop and block nested loop algorithms including:
  - **Scan inner loop forward and backward alternately, to make use of the blocks remaining in buffer (with LRU replacement)**
  - **Use index on inner relation if available (next slide)**

E.g. Compute  $student \bowtie takes$ , with  $R=student$  as the outer relation,  $S=takes$  the inner relation:

- Number of records for *student* : 5,000, and for *takes*: 10,000
- Number of blocks for *student*: 100 and for *takes*: 400

# Index Join (IJ)

- (1) Create an index for S.C if one is not already there
- (2) For each  $r \in R$  find matching records in S  
( $s.C=r.C$ ) and output (r,s)



# Example of Nested-Loop Join Costs

- Compute  $student \bowtie takes$ , with *student* as the outer relation.
  - Number of records of *student*: 5,000      *takes*: 10,000
  - Number of blocks of *student*: 100      *takes*: 400
- Let *takes* have a primary B<sup>+</sup>-tree index on the attribute *ID*, which contains 20 entries in each index node. Since *takes* has 10,000 tuples, the height of the tree is 4.
- 4 nodes accessed from the index +1 more access to the actual data
- *student* has 5000 tuples
- Cost of indexed nested loops join
  - $100 + 5000 * 5 = 25,100$  block transfers
- Cost of block nested loops join
  - $400 * 100 + 100 = 40,100$  block transfers assuming worst case memory
    - may be significantly less with more memory

# Hash Join (HJ)

- Hash function  $h(v)$ , range  $1 \rightarrow k$


## Algorithm

(1) Hashing stage (bucketizing): hash tuples into buckets

- Hash R tuples into  $G_1, \dots, G_k$  buckets
- Hash S tuples into  $H_1, \dots, H_k$  buckets

(2) Join stage: join tuples in matching buckets



- For  $i = 1$  to  $k$  do  
    match tuples in  $G_i, H_i$  buckets





# Hash Join (HJ)

- $H(k) = k \bmod 3$




40	T1
60	T2
30	T3
10	T4
20	T5


10	T6
60	T7
40	T8
20	T9

# Hash Equi Join (HJ) of G and H

- Step (1): Hashing stage




- Step (2): Join stage



# HJ: Critical for Parallel Joins


- Step (1): Hashing stage for G



–

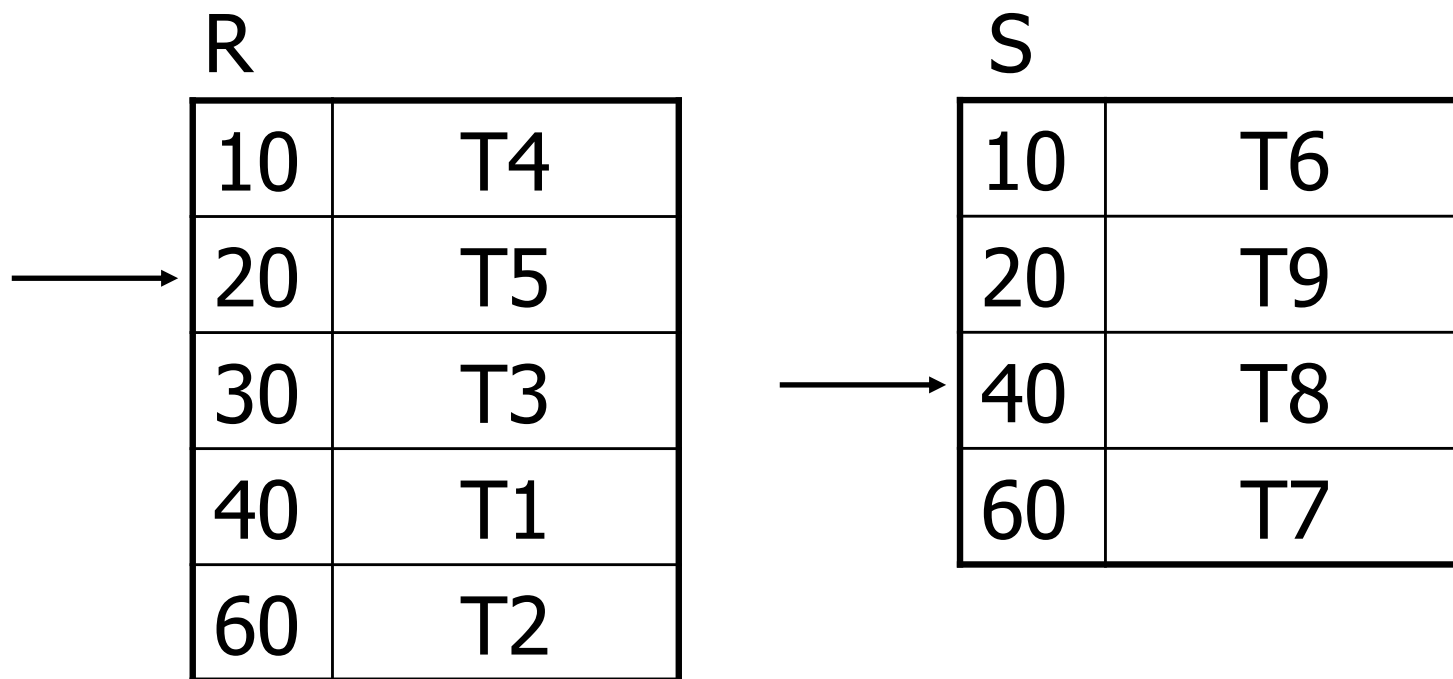
– Ditto for H: H1, H2, H3 ...

- Step (2): Join stage



# Sort-Merge Join (SMJ)

- Sort the relations first and then join



Quicksort can be used if data fits in memory  
External sor is needed otherwise.

# Sort-Merge Join (SMJ)

(1) if R and S not sorted, sort them

(2)  $i \leftarrow 1; j \leftarrow 1;$

While  $(i \leq |R|) \wedge (j \leq |S|)$  do

if  $R[i].C = S[j].C$  then outputTuples

else if  $R[i].C > S[j].C$  then  $j \leftarrow j+1$

else if  $R[i].C < S[j].C$  then  $i \leftarrow i+1$

R  
i

10	T4
20	T5
30	T3
40	T1
60	T2

S  
j

10	T6
20	T9
40	T8
60	T7

(no incrementr to I or J when we have equality?)

# Sort-Merge Join (SMJ)

(1) if R and S not sorted, sort them

(2)  $i \leftarrow 1; j \leftarrow 1;$

While  $(i \leq |R|) \wedge (j \leq |S|)$  do

if  $R[i].C = S[j].C$  then {outputTuples;  
 $j \leftarrow j+1; i \leftarrow i+1$  }

else if  $R[i].C > S[j].C$  then  $j \leftarrow j+1$

else if  $R[i].C < S[j].C$  then  $i \leftarrow i+1$

R  
i

10	T4
20	T5
30	T3
40	T1
60	T2

S  
j

10	T6
20	T9
40	T8
60	T7

This only works when the keys do not contain duplicates

# External Merge Sort

❖ *Idea: Divide and conquer:* sort subfiles and merge into larger sorts

Pass 0 → Only one memory block is needed

Pass  $I > 0$  → Only three memory blocks are needed

