

CS143: Homework 6

Problem A: T1 is the oldest transaction, and T3 is the youngest.

T1	T2	T3
write D		
	write C	
read C		
		read C
		write A
	read A	
write A		

1. Is this first schedule conflict-serializable¹?

Answer: We have that $T2 \rightarrow T3$ on C and $T3 \rightarrow T2$ on A. The directed cycle $T3 \xrightarrow{C} T2 \xrightarrow{A} T3$ tells us that the schedule is not conflict-serializable.

Now assume that the transaction manager uses a 2PL protocol where each exclusive/shared lock is set just before it is needed for the write/read action.

2. If we do not use any deadlock prevention strategy, will the resulting transactions (i) complete, or (ii) deadlock¹? If your answer is (i), show a completed schedule;

Answer: if it is (ii) show the schedule up to the deadlock. So T2 sets a lock-X C just before write C, and because of 2PL it must keep until it is done with its read A. Thus when T1 does lock-S C it stops with a wait-for arc $T1 \rightarrow T2$. Ditto for T3 that sets the wait-for arc $T3 \rightarrow T2$. There is no cycle: thus no deadlock. T2 completes. Then T3 and T1 will complete also.

3. If we use a wait-die deadlock prevention strategy will the resulting transactions (i) complete, or (ii) deadlock¹? If your answer is (i), show a completed schedule; if it is (ii) show the schedule up to the deadlock.

Answer: Obviously there is going to be no deadlock since Wait-Die prevents deadlocks. Here, T1 will wait for the older transaction to release C. But T3 is younger than T2; so when it requests C it will actually die. Again a correct completion sequence is T2, T1, T3.

From now on consider the second schedule below and use that in your answer:

T1	T2	T3
write D		
	write C	
read C		
		write A
		read C
	read A	
write A		

¹Justify your answer using the applicable graph

4. Is this second schedule conflict-serializable¹?

Answer: Here too, we have that $T2 \rightarrow T3$ on C and $T3 \rightarrow T2$ on A. This cycle tells us that the schedule is not conflict-serializable.

Now assume that the transaction manager uses a 2PL protocol where each exclusive/shared lock is set just before it is needed for the write/read action, and answer the following questions for this second schedule.

5. If we do not use any deadlock prevention strategy, will the resulting transactions (i) complete, or (ii) deadlock¹? If your answer is (i), show a completed schedule; if it is (ii) show the schedule up to the deadlock.

Answer: So T2 sets a lock-X C just before write C, and because of 2PL it must keep until it is done with its read A. Thus when T1 does lock-S C it stops with a wait-for arc $T1 \rightarrow T2$. Next, T3 does lock-X A, and then by requesting lock-s C it sets a wait-for arc $T3 \rightarrow T2$. Thus, so far, no cycles and no deadlock. But then T2 does lock-S A and sets a wait-for arc: $T2 \rightarrow T3$. Thus we now have deadlock as per the directed cycle $T3 \xrightarrow{\text{wait-for}} T2 \xrightarrow{\text{wait-for}} T3$ in the wait-for graph.

6. If we use a Wound-Wait deadlock prevention strategy will the resulting transactions (i) complete, or (ii) deadlock¹? If your answer is (i), show a completed schedule; if it is (ii) show the schedule up to the deadlock.

Answer: Obviously there is going to be no deadlock since Wound-Wait prevents deadlocks. Here, T1 is the older transaction requesting C held by the younger T2. Thus T1 wounds T2 which dies releasing C. Then T3 locks A and C and completes releasing them both. Finally T2 restarts and completes. Thus the correct completion sequence is T3, T1 and finally T2.

Problem B 1. Consider the following schedule: (w3(A) means that transaction T3 writes A, C3: T3 commits):

w3(A) r1(A) c3 w1(B) c1 r2(B) w2(C) r4(B) c2c4

- (a) Is it a serial schedule?
- (b) Is the schedule conflict serializable? If so, what are all the equivalent serial schedules?
- (c) Is the schedule recoverable? If not, can we make it recoverable by moving a single commit operation to a different position?
- (d) Is the schedule cascadeless? If not, can we make it cascadeless by moving a single commit operation to a different position?

T1	T2	T3	T4
		w(A)	
r(A)		commit	
w1(B)			
commit			
	r(B)		
	w(C)		
	commit		
			r(B)
			commit

- (a) this is not a serial schedule
- (b) $T3 \rightarrow T1$ (due to A) $T1 \rightarrow T2$ (do to B) $T1 \rightarrow T4$ (due to B)

No cycle. The schedule is serializable. e.g. T3 before T1 before T2 and before T4

(c) Dirty data with current schedule: A and B. A is written by T3 and read by T1 after T1 commits. Then, the recovery process will undo w1(B) and generate a schedule that could not have been generated by an equivalent serial schedule (unrecoverability of schedules).

d) Dirty data with current schedule: A and B. A is written by T3 and read by T1 after T1 commits—no cascading of rollback is needed because of A. However, B is read by T1 and T4 before they commit: thus we have a cascade of rollbacks if T1 fails. To make this schedule recoverable and cascadeless, we must delay the commit of T1 until T2 and T4 commit.