

SQL Notes—Part 2

SELECT ... FROM ... WHERE can express projection, joins, and selection

But in addition to those SQL also provides specialized constructs to support:

1. Aggregates
2. Explicit Joins, and
3. Set Union, and
4. Set Difference from Relational Algebra.

Explicit Join

For every person, return the person, their father and their mother

```
select fatherChild.child, father, mother  
from   motherChild, fatherChild  
where  fatherChild.child = motherChild.child
```

```
select fatherChild.child, father, mother  
from   motherChild join fatherChild on  
       fatherChild.child = motherChild.child
```

UNION

Within a `select` statement one cannot express unions.

An explicit construct is needed:

```
select ...
union [all]
select ...
```

With `union`, duplicates are eliminated (also those originating from projection).

With `union all` duplicates are kept.

Positional Notation of Attributes

```
select father, child  
from   fatherChild  
union  
select mother, child  
from   motherChild
```

Which are the attribute names of the result?
Those of the first operand!

- SQL matches attributes in the same position
- SQL renames the attributes of the second operand

union all will keep duplicates

MotherChild

| mother | child |
|--------|-------|
| Lisa | Mary |
| Lisa | Greg |
| Anne | Kim |
| Anne | Phil |
| Mary | Andy |
| Mary | Rob |

FatherChild

| father | child |
|--------|-------|
| Steve | Frank |
| Greg | Kim |
| Greg | Phil |
| Frank | Andy |
| Frank | Rob |

Result of
the Union

father child

| | |
|-------|-------|
| Greg | Frank |
| Greg | Kim |
| Greg | Phil |
| Frank | Andy |
| Frank | Rob |
| Lisa | Mary |
| Lisa | Greg |
| Anne | Kim |
| Anne | Phil |
| Mary | Andy |
| Mary | Rob |

Correct (if we want to treat fathers and mothers as parents):

```
select father as parent, child  
from   fatherChild  
union  
select mother as parent, child  
from   motherChild
```

For set difference: Replace ‘union’ by ‘except’

Set Difference: except

Find children who have a father but not a mother:

select child as childname

from FatherChild

except

select child from Mothedhild

MotherChild

| mother | child |
|--------|-------|
| Lisa | Mary |
| Lisa | Greg |
| Anne | Kim |
| Anne | Phil |
| Mary | Andy |
| Mary | Rob |

FatherChild

| father | child |
|--------|-------|
| Steve | Frank |
| Greg | Kim |
| Greg | Phil |
| Frank | Andy |
| Frank | Rob |

Person

| name | age | income |
|-------|-----|--------|
| Andy | 27 | 21 |
| Rob | 25 | 15 |
| Mary | 55 | 42 |
| Anne | 50 | 35 |
| Phil | 26 | 30 |
| Greg | 50 | 40 |
| Frank | 60 | 20 |
| Kim | 30 | 41 |
| Mike | 85 | 35 |
| Lisa | 75 | 87 |

Queries from Hw

customer(customer-name, street, city)

branch(branch-name, city)

account(customer-name, branch-name, account-number)

- (c) *Find the branches that do not have any accounts.*
- (d) *Find the customer names who do not have any account in the ‘Region12’ branch.*
- (e) *Find the customer names who have accounts in all the branches located in ‘Los Angeles’.*
- (f) *Find the customer names who have only one account.*

To express these in SQL, we need nested subqueries

MotherChild

| mother | child |
|--------|-------|
| Lisa | Mary |
| Lisa | Greg |
| Anne | Kim |
| Anne | Phil |
| Mary | Andy |
| Mary | Rob |

FatherChild

| father | child |
|--------|-------|
| Steve | Frank |
| Greg | Kim |
| Greg | Phil |
| Frank | Andy |
| Frank | Rob |

Person

| name | age | income |
|-------|-----|--------|
| Andy | 27 | 21 |
| Rob | 25 | 15 |
| Mary | 55 | 42 |
| Anne | 50 | 35 |
| Phil | 26 | 30 |
| Greg | 50 | 40 |
| Frank | 60 | 20 |
| Kim | 30 | 41 |
| Mike | 85 | 35 |
| Lisa | 75 | 87 |

Nested Queries (Example)

“Name and income of Frank’s father”

```
select f.name, f.income  
from person f, fatherChild fc  
where f.name = fc.father and fc.child = 'Frank'
```

```
select f.name, f.income  
from person f  
where f.name = (select fc.father  
                 from fatherChild fc  
                 where fc.child = 'Frank' )
```

Instead of = (i.e., equal to any) we could have used IN (i.e. in the set)

Nested Queries: Example

Name and income of the fathers of persons who earn more than 20k.

```
select distinct f.name, f.income  
from person f, fatherChild fc, person c  
where f.name = fc.father and  
fc.child = c.name and c.income > 20
```

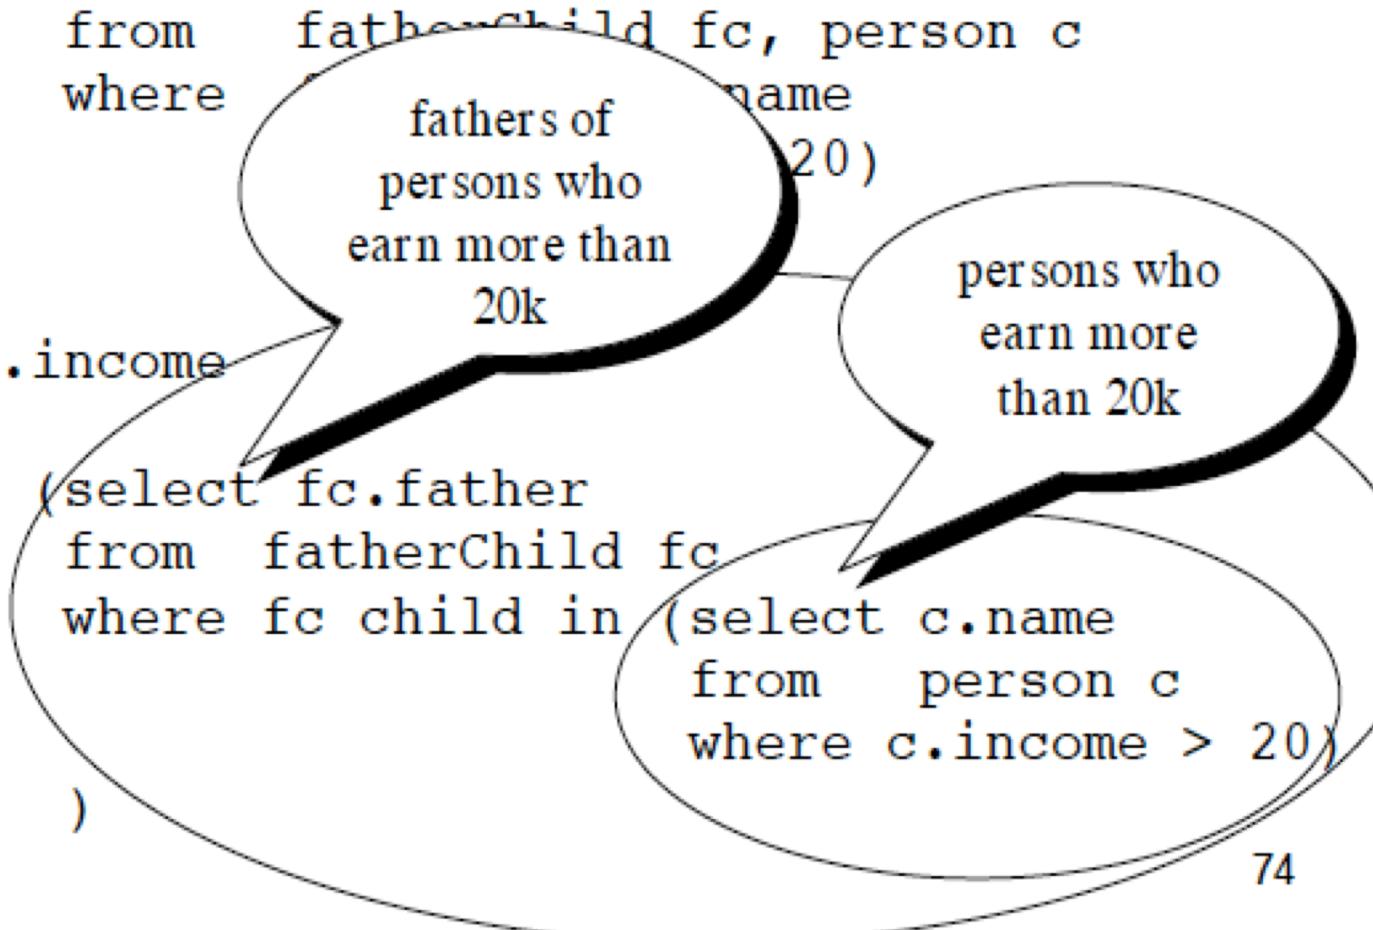
```
select f.name, f.income  
from person f  
where f.name = any
```

```
(select fc.father  
from fatherChild fc, person c  
where fc.child = c.name and  
c.income > 20)
```

fathers of persons
who earn more
than 20k

Name and income of the fathers of persons who earn more than 20k.

```
select f.name, f.income  
from person f  
where f.name in (select fc.father  
                  from fatherChild fc, person c  
                  where fc.child = c.name  
                        and c.income > 20)  
  
select f.name, f.income  
from person f  
where f.name in (select fc.father  
                  from fatherChild fc  
                  where fc.child in (select c.name  
                                     from person c  
                                     where c.income > 20))
```



Using nested subqueries to express joins does not bring obvious benefits, in general.

However nesting combined with constructs such as ALL, NOT EXIST, and IN allows us to express queries that cannot be expressed without these new constructs.

Nested Queries: Example with all

“Persons who have an income that is higher than the income of all persons younger than 30”

```
select name  
from person  
where income >= all (select income  
                      from person  
                      where age < 30 )
```

“Name and age of the mothers all of whose children are at least 18”

```
select m.name, m.age
from   person m join motherChild mc
          on m.name = mc.mother
where  18 <= all (select c0.age
                  from   motherChild mc0 join person c0
                  on mc0.mother = c0.name
                  where  mc0.mother = mc.mother)
```

Nested Queries: Example with exists

An expression with the operator `exists` is true if the result of the subquery is not empty.

Example: “Persons with at least one child”

```
select *
from person p
where exists (select *
               from fatherChild fc
               where fc.father = p.name)
      or
exists (select *
               from motherChild mc
               where mc.mother = p.name)
```

Note: the attribute `name` refers to the table in the outer `from` clause.

All the queries with nesting in the previous examples are equivalent to some unnested query. So, what's the point of nesting?

Example: “Persons without a child”

```
select *
from   person p
where  not exists (select *
                     from   fatherChild fc
                     where  fc.father = p.name)
and
not exists (select *
             from   motherChild mc
             where  mc.mother = p.name)
```

This cannot be expressed equivalently as a “select from where” query.

Why?

Solution with `not exists`

“Name and age of the mothers all of whose children are at least 18”

```
select m.name, m.age
from   person m join motherChild mc
          on m.name = mc.mother
where  not exists
       (select *
        from   motherChild mc0 join person c0
          on mc0.mother = c0.name
        where  mc0.mother = mc.mother and
               c0.age < 18)
```

What about childless mothers? Note difference w.r.t. natural language.

Replace `mc0.mother=c0.name` by `mc0.chkid=c0.name`

Exercise 9

“Name and age of the mothers all of whose children are at least 18”

Approach 1: Subquery with `all`

Approach 2: Subquery with `min`

Approach 3: Subquery with `not exists`

Equivalent formulation with max

“Persons who have an income that is higher than the income of all persons younger than 30”

```
select name  
from person  
where income >= (select max(income)  
                   from person  
                   where age < 30)
```

Is this always equivalent to the previous query that used ALL?

Name and age of mothers whose children are all ≥ 18 : using min

```
select m.name, m.age
from   person m join motherChild mc
          on m.name = mc.mother
where  18 <= (select min(c0.age)
               from   motherChild mc0 join person c0
               on mc0.mother = c0.name
               where mc0.mother = mc.mother)
```

The EXISTS construct is also supported... but it is not as useful as NOT EXISTS

An expression with the operator `exists` is true if the result of the subquery is not empty.

Example: “Persons with at least one child”

```
select *
from person p
where exists (select *
               from fatherChild fc
               where fc.father = p.name)
or
exists (select *
        from motherChild mc
        where mc.mother = p.name)
```

Note: the attribute name refers to the table in the outer `from` clause.

Nested Queries: Visibility

persons having at least one child.

```
select *
from person
where exists (select *
               from fatherChild
               where father = name)
      or
exists (select *
        from motherChild
        where mother = name)
```

The attribute **name** refers to the table **person** in the outer **from** clause.

Nested Queries: Comments

- Visibility rules:
 - it is not possible to refer to a variable defined in a block ~~Inside~~^{below} the current block
 - if an attribute name is not qualified with a variable or table name, it is assumed that it refers to the “closest” variable or table with that attribute
- In each block, one can refer to variables defined in the same block or in surrounding blocks
- Semantics: the inner query is executed for every tuple of the outer query

More on Visibility= Variable Scope

Note: This query is incorrect:

```
select *
from employee
where dept in (select name
                from department D1
                where name = 'Production')
      or
dept in (select name
          from department D2
          where D2.city = D1.city)
```

| | | | |
|----------|------|----------|------|
| employee | name | lastName | dept |
|----------|------|----------|------|

| | | | |
|------------|------|---------|------|
| department | name | address | city |
|------------|------|---------|------|

SET DIFF and NESTING

employee(name, lastname, dept)

Can one express set difference by way of nesting?

```
select name from employee  
except  
select lastName as name from employee
```

```
select name  
from employee  
where name not in (select lastName  
                     from employee)
```

“The person (or the persons) that have the highest income”

```
select *
from person
where income = (select max(income)
                  from person)
```

Or:

```
select *
from person
where income >= all (select income
                      from person)
```

Can we do it with not exist?

Nested Queries: Conditions on Several Attributes

The persons which have a unique combination of age and income

(that is, persons for whom the pair (age, income) is different from the corresponding pairs of all other persons).

```
select *
from person p
where (age,income) not in
      (select age, income
       from person
       where name <> p.name)
```

MotherChild

| mother | child |
|--------|-------|
| Lisa | Mary |
| Lisa | Greg |
| Anne | Kim |
| Anne | Phil |
| Mary | Andy |
| Mary | Rob |

FatherChild

| father | child |
|--------|-------|
| Steve | Frank |
| Greg | Kim |
| Greg | Phil |
| Frank | Andy |
| Frank | Rob |

Person

| name | age | income |
|-------|-----|--------|
| Andy | 27 | 21 |
| Rob | 25 | 15 |
| Mary | 55 | 42 |
| Anne | 50 | 35 |
| Phil | 26 | 30 |
| Greg | 50 | 40 |
| Frank | 60 | 20 |
| Kim | 30 | 41 |
| Mike | 85 | 35 |
| Lisa | 75 | 87 |

Explicit Join

For every person, return the person, their father and their mother

```
select fatherChild.child, father, mother  
from   motherChild, fatherChild  
where  fatherChild.child = motherChild.child
```

```
select fatherChild.child, father, mother  
from   motherChild join fatherChild on  
       fatherChild.child = motherChild.child
```

MotherChild

| mother | child |
|--------|-------|
| Lisa | Mary |
| Lisa | Greg |
| Anne | Kim |
| Anne | Phil |
| Mary | Andy |
| Mary | Rob |

FatherChild

| father | child |
|--------|-------|
| Steve | Frank |
| Greg | Kim |
| Greg | Phil |
| Frank | Andy |
| Frank | Rob |

Joins and Information Preservation by Nulls

Find the Mother and Father of each child

MotherChild

| mother | child |
|--------|-------|
| Lisa | Mary |
| Lisa | Greg |
| Anne | Kim |
| Anne | Phil |
| Mary | Andy |
| Mary | Rob |

FatherChild

| father | child |
|--------|-------|
| Steve | Frank |
| Greg | Kim |
| Greg | Phil |
| Frank | Andy |
| Frank | Rob |

MotherChild \bowtie *FatherChild*

*With join: Unmatched values
are lost from both
left (i.e., motherChild) and
right (i.e., fatherChild)*

*But with we can use
NULL values to fill in for
the missing value:
outerjoins*

(Lisa, Mary, Null) *left-outer join*

....

(Null, Steve, Frank) *left-outer join*

Outer Join: Examples

```
select fatherChild.child, father, mother  
from   motherChild join fatherChild  
       on motherChild.child = fatherChild.child
```

Unmatched values are lost from both left (i.e., motherChild) and right (i.e., fatherChild)

```
select fatherChild.child, father, mother  
from   motherChild left outer join fatherChild  
       on motherChild.child = fatherChild.child
```

Values on the left are preserved, since unmatched ones are padded with null values-

```
select fatherChild.child, father, mother  
from   motherChild right outer join fatherChild  
       on motherChild.child = fatherChild.child
```

Values on the right are preserved, since unmatched ones are padded with null values

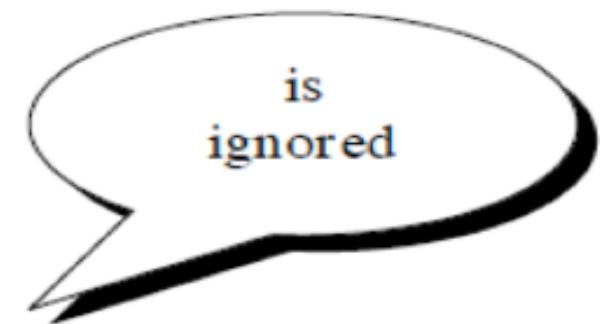
```
select fatherChild.child, father, mother  
from   motherChild full outer join fatherChild  
       on motherChild.child = fatherChild.child
```

Values on both right and Left are preserved, since unmatched ones are padded with null values

Null Values

Select name from person where income <100

| person | name | age | income |
|--------|------|-----|--------|
| | Andy | 27 | 30 |
| | Rob | 25 | NULL |
| | Mary | 55 | 36 |
| | Anne | 50 | 36 |

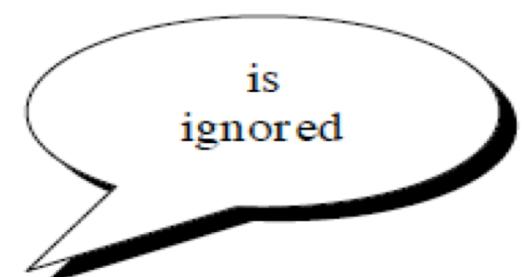


Three-valued logic: True, False and Logical Null

Aggregate Operators and Null Values

```
select avg(income) as meanIncome  
from person
```

| person | name | age | income |
|--------|------|-----|--------|
| | Andy | 27 | 30 |
| | Rob | 25 | NULL |
| | Mary | 55 | 36 |
| | Anne | 50 | 36 |



is
ignored

| meanIncome |
|------------|
| 34 |

count and Null Values

```
select count(*)  
from person
```

Result = number of tuples
= 4

```
select count(income)  
from person
```

Result = number of values
different from NULL
= 3

```
select count(distinct income)  
from person
```

Result = number of distinct
values (excluding
NULL)
= 2

| person | name | age | income |
|--------|------|-----|--------|
| | Andy | 27 | 21 |
| | Rob | 25 | NULL |
| | Mary | 55 | 21 |
| | Anne | 50 | 35 |

SQL Queries from Hw by outerjoins

Customer(customer-name, street, city)

Branch(branch-name, city)

Account(customer-name, branch-name, account-number)

%the join of Customer and Account has the following attributes

(customer-name, street, city, customer-name, branch-name, account-number)

c) Find the branches that do not have any account.

SELECT customer-name

FROM Customer LEFT OUTERJOIN Account ON

Customer.customer-name=Account.customer-name

WHERE ISNULL(account-number)

Queries from Hw by Outerjoins

Customer(customer-name, street, city)

Branch(branch-name, city)

Account(customer-name, branch-name, account-number)

(c) Find the branches that do not have any accounts.

```
SELECT    customer-name  
FROM      Customer LEFT OUTERJOIN Account  
          ON Customer.customer-mame=Account.customer-name  
HAVING    COUNT(account-number)= 0
```

But COUNT(*) will not work because this counts the number of tuple occurrences, independent of their values.