

ISAM Indexes

Topics to Learn

- Important concepts
 - Dense index vs. sparse index
 - Primary index vs. secondary index
(= clustering index vs. non-clustering index)
 - Tree-based vs. hash-based index
- Tree-based index
 - Indexed sequential file
 - B+-tree
- Hash-based index
 - Static hashing
 - Extendible hashing

Basic Problem

- `SELECT *`
`FROM Student`
`WHERE sid = 40`

sid	name	GPA
20	Elaine	3.2
70	Peter	2.6
40	Susan	3.7

- How can we answer the query?

Random-Order File

- How do we find sid=40?

sid	name	GPA
20	Susan	3.5
60	James	1.7
70	Peter	2.6
40	Elaine	3.9
30	Christy	2.9

Sequential File

- Table sequenced by sid. Find sid=40?


sid	name	GPA
20	Susan	3.5
30	James	1.7
40	Peter	2.6
50	Elaine	3.9
60	Christy	2.9

Binary Search

- 100,000 records
- Q: How many blocks to read?
- Any better way?
 - In a library, how do we find a book?


Basic Idea

- Build an “index” on the table
 - An auxiliary structure to help us locate a record given a “key”



Dense, Primary Index

Dense Index Sequential File




- Primary index (clustering index)
 - Index on the search key
- Dense index
 - (key, pointer) pair for every record
- Find the key from index and follow pointer
 - Maybe through binary search
- Q: Why dense index?
 - Isn't binary search on the file the same?

Why Dense Index?

- Example
 - 10,000,000 records (900-bytes/rec)
 - 4-byte search key, 4-byte pointer
 - 4096-byte block. Unspanned tuples
- Q: How many blocks for table (how big)?
 - How many unspanned records per block: $4096 \div 900 = 4$
 - How many blocks $10,000,00 \div 4 = 2,500,000$.
- Q: How many blocks for index (how big)?
 - For dense index, one search key and one pointer per record requiring $4+4=8$ bytes
 - $8 \times 10,000,000 = 80,000,000$ bytes needed for dense index
 - How many keys per block: $4096 \div 8 = 512$
 - How many blocks for indices: $80000000 \div 512 = 156,250$ Thus 156250 are needed for the index.


Sparse, Primary Index

Sparse Index



- Sparse index
 - (key, pointer) pair per every "block"
 - (key, pointer) pair points to the first record in the block
- Q: How can we find 60?

Multi-level index



Q: Why multi-level index?

Q: Does dense, 2nd level index make sense?

Secondary (non-clustering) Index

Sequence
field



30	
50	

20	
70	


80	
40	

100	
10	


90	
60	

- Secondary (non-clustering) index
 - When tuples in the table are not ordered by the index search key
 - Index on a non-search-key for sequential file
 - Unordered file
- Q: What index?
 - Does sparse index make sense?

Sparse and secondary index?



Secondary index




- First level is typically dense
- Sparse from the second level

Important terms

- Dense index vs. sparse index
- Primary index vs. secondary index
 - Clustering index vs. non-clustering index
- Multi-level index
- Indexed sequential file
 - Sometimes called ISAM (indexed sequential access method)
- Search key (\neq primary key)

Insertion

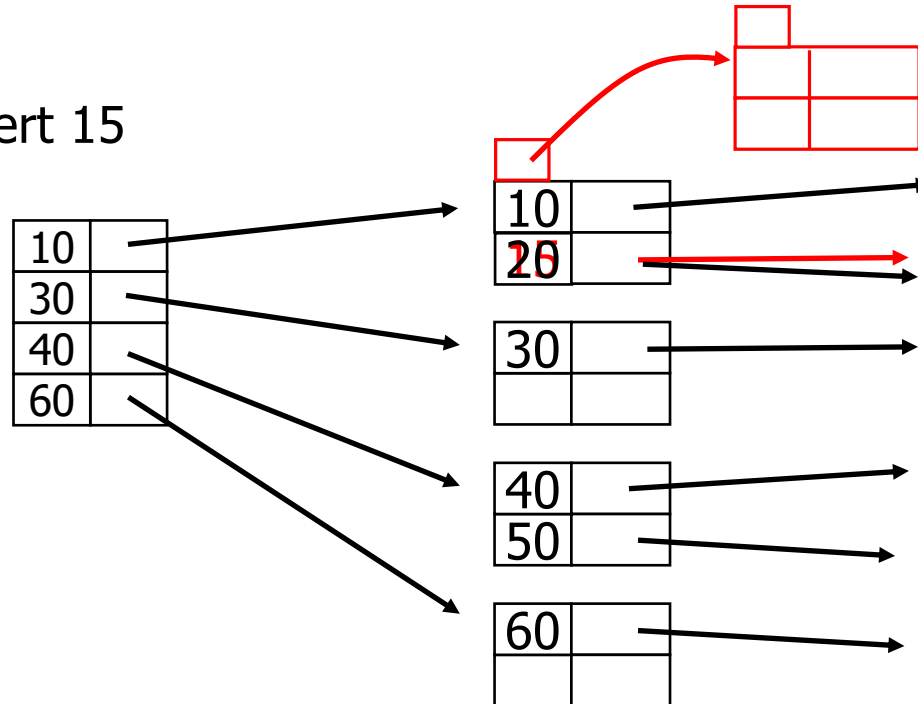
Insert 35



Q: Do we need to update higher-level index?

Insertion


Insert 15



Q: Do we need to update higher-level index?

Insertion

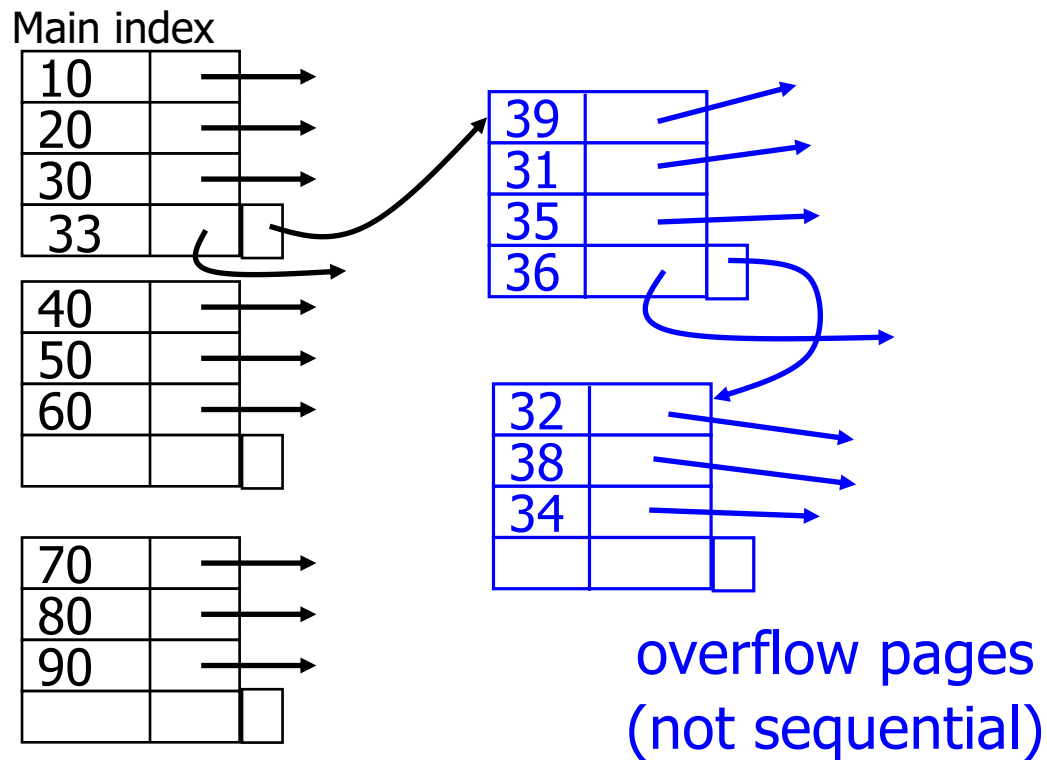
Insert 15



Q: Do we need to update higher-level index?

Potential performance problem

After many insertions...



Traditional Index (ISAM)

- Advantage
 - Simple
 - Sequential blocks
- Disadvantage
 - Not suitable for updates
 - Becomes ugly (loses sequentiality and balance) over time